Introduction to programming for astronomers (fall 2015)

Thijs Coenen

October 14, 2015

Contents

Chapter 1

Introduction

In astronomy, and outside of it, programming is becoming an increasingly important skill. The general availability of lots of compute power creates new opportunities for astronomy. There are three general areas in astronomy where computers play an increasingly important role. First, the latest generations of observatories produce high data rates that at times need to be searched in real time for interesting signals. Second, there is a move to more open data sharing and public archiving of observational data, which creates opportunities for data mining. Third the availability of massive amounts of computational power allows increasingly detailed astrophysical simulations to be performed.

1.1 Basic Linux and Coding for Astronomy & Astrophysics

The aim of the course *Basic Linux and Coding for AA (5214BLCF3Y)* is to to get you up and running in Linux with some basic knowledge of its command line interface and teach you the basics of programming. This course will use the Python programming language (more specifically Python 2.7 which, despite Python 3 being available, is still the most prevalent in research). Python was chosen for its relatively straightforward syntax, its free availability, the wide range of libraries in its "eco system". For the fields physics and astronomy in particular, Python provides a large number of libraries for scientific computing and is displacing packages like IDL[1] and Matlab[2].
The Linux part of our course will teach you some of the backgrounds and commands that you will need to be productive. More specifically:

- You should attain basic knowledge of the UNIX/Linux file system layout and the use of permissions on UNIX/Linux to control file access.

- You should be able to access the command line (we will be using the BASH shell throughout this course).

- You should be able to manage running processes on a Linux system, to start or stop programs and to find out which processes are running.

- You should be able to navigate the UNIX/Linux file system and find files and programs.

- You should be able to manipulate text files through the command line and search for specific content.

The programming part of this course aims to teach you the basics of programming in general and programming in Python in particular. Because of the limited extent of this course, we will focus our attention on a practical subset of the Python language — and test whether you attain working knowledge of that subset. Specifically:

- You should be able to design and implement a properly structured Python program written in the standard Python style. You will also learn to properly document you program, use proper variable and function naming, how to modularize you program using functions, and how to make your program reusable.

---

[1]`http://www.exelisvis.com/ProductsServices/IDL.aspx`
[2]`http://www.mathworks.com/products/matlab/`

- You should be able to read and write simple text data files.

- You should be able create basic plots using Matplotlib.

- You should be familiar with Numpy and Scipy and be able to use these libraries to implement simple scientific calculations efficiently.

A further more general goal is that you develop problem solving skills using programming tools.

- Learn how to find help when programming. That includes how to read the documentation available with Python and its libraries, and which resources are available on-line.

- Develop the skills to, when presented with new data files, inspect that data, create plots and write small tools to work with them.

1.2   Conventions

This reader uses a number of typographical conventions differentiate normal text, programming code snippets, commands you enter and program output. In running text you may come across monospaced examples or names. When these examples are underlined, `like this`, they are (partial) commands or snippets you may enter. When programs require key presses that are not really commands (not followed by an enter), they are type set as for instance ⌷q⌷ or ⌷control⌷ + ⌷z⌷ if two keys must be pressed at the same time. Program output and names of programs are typeset `like this`. Examples of shell usage, interactive Python sessions and Python source code are below. First a shell session:

```
Gretchen:~ thijscoenen$ cd coolstuff/
Gretchen:coolstuff thijscoenen$ ls
Whatever        alice.txt       blah.txt        hello               noperm
Gretchen:coolstuff thijscoenen$ head alice.txt
ALICE'S ADVENTURES IN WONDERLAND

Lewis Carroll


CHAPTER I. Down the Rabbit-Hole

Alice was beginning to get very tired of sitting by her sister on the
bank, and of having nothing to do: once or twice she had peeped into the
book her sister was reading, but it had no pictures or conversations in
Gretchen:coolstuff thijscoenen$
```

The second example is an interactive Python session:

```
>>> l = range(10)
>>> print l[:5]
[0, 1, 2, 3, 4]
>>> 1 / 0
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ZeroDivisionError: integer division or modulo by zero
>>>
```

Here the prompt >>> is followed on the same line by some Python code and followed by some output from that command on the following line(s). The last example is Python source code:

```python
#!/usr/env/bin python
'''
Print a custom greeting.
'''


def say_hello(person):
    print 'Hello {}'.format(person)


if __name__ == '__main__':
    say_hello('Jan Janssen')
```

1.3   Contents of this reader

This reader consists of 3 parts. The first part explains the basics of the Linux command line, the basics of Python and some libraries useful for scientific computing. The second part contains a number of assignments that will be used during the tutorials. Finally, the third part (the appendix) contains some extra background information.

Part I

Reading

Chapter 2

Linux and its Command Line

In this chapter I will quickly introduce you to the very basics of using Linux through its command line interface. After a little history of Linux and operating systems like it, I will explain how you get access to the command line and go over some of the commands you will need. This chapter can be skipped safely if you are already familiar with the Linux command line. The examples in this chapter use the so called Bourne Again Shell (BASH). This shell comes standard with many flavors of Linux or UNIX.

2.1   UNIX and Linux

UNIX is an operating system developed in the early 1970s at AT&T Bell Labs. Unlike some of its contemporaries UNIX was designed as a multi user system from the start, it used a hierarchical file system (see 2.1.2) and consisted of many small programs that could be combined to perform complex operations. UNIX became popular in academia, where UNIX like systems are still used. In the early 1990s Linus Torvalds developed a new UNIX like operating system that he could run on his own PC. This operating system became known as Linux and, while initially developed by volunteers and hobbyists, it was soon picked up by businesses. Nowadays a large fraction of Linux' development is done by very large computer companies like Google, IBM and Samsung. Apple's Mac OS X is another popular UNIX version.

Today Linux is used on computers ranging from mobile phones (e.g. Google's Android) to the largest super computer clusters. It furthermore runs a large fraction of the Internet's infrastructure (e.g. routers and web servers). Because UNIX was developed before graphical displays became generally available, its earliest interfaces were all text driven. Although graphical user interfaces are available for UNIX and Linux nowadays the textual interface persists. This textual user interface is also called a *command line* interface and on UNIX is provided through a so-called *shell*. Several different shells are available, but this chapter will only explain the basics of the Bourne Again Shell (BASH).

2.1.1   Users and Permissions

Because UNIX was designed as a multi user operating system, it has a notion of users. Associated with every user are a level of access to the operating system. The so-called root user has full access to a UNIX system while the other users have more restricted access. On your own computer you will likely be able to log in as root while on shared computers you will only have access to your own files. Do note that even if you can log in as root, it is absolutely a bad idea to do so for day-to-day work. Since the root is all-powerful, a mistake while logged in as root can be much worse than a mistake while logged in as an ordinary user. Furthermore the software that you use may contain flaws or security problems that become a problem when that software (running as root) has full access to the computer. Furthermore each user is also assigned to a group of users that have the same system privileges.

UNIX controls access to files and directories based on ownership and so-called *permissions*. Each file and directory on a UNIX system has an *owner* (one of the users on that system) and permissions. UNIX has three permissions: *read*, *write* and *execute*. Read permission is needed to be able to view the contents of that file. Write permission is needed on directories to create files in that directory. For files write permission is needed to change or erase them. Execute permission is needed on directories to be able to view their contents and on files to execute them.

Permission may also be set at the *group* level or for all *other* users with access to a computer system (see Section 2.2.12).

2.1.2 The File System

UNIX uses a hierarchical file system meaning that directories in it can be arbitrarily nested. Each directory can contain files or sub directories — the former directory is also referred to as *parent* while the latter are referred to as *children*. The directory hierarchy is a tree with the directory at the base of that tree referred to as *root*, denoted as /. When specifying locations, or *paths*, in a file system two different notations may be used. *Absolute paths*, on the one hand, describe the location of a file or directory in absolute terms, i.e. without reference to some other location on the file system. *Relative paths*, on the other hand, specify a location in the file system relative to some other location in the file system (usually the current location of the user). Absolute paths start at the root, i.e. they start with /, while relative paths do not.

Because a UNIX system may be shared between many users, each user has his or her own directory to store files in. This special directory is called a *home directory*. Home directories can be found in /home on Linux systems and in /Users for Mac OS X systems. E.g. my user (thijscoenen) has the home directory /Users/thijscoenen on a Mac OS X system and /home/thijscoenen on a Linux system.

The UNIX(-like) operating systems have very specific file system layouts, with some directories assigned to programs or even parts of the operating system itself. Most directories that reside directly in the root directory are in fact system directories and as a normal user you will rarely need access to them. Below, as an example, the contents of my Mac OS X machine's root directory none of which is part of my own files or directories[1]

```
Gretchen:coolstuff thijscoenen$ ls /
Applications    System        cores        mach_kernel    tmp
Developer       Users         dev          net            usr
Library         Volumes       etc          private        var
Network         bin           home         sbin
Gretchen:coolstuff thijscoenen$
```

UNIX also allows so-called *hidden files* that are normally invisible. Any filename or directory name that starts with a dot will be hidden. Hidden files are generally used for settings or for those files that a user does not need access to directly (only through some program). E.g. the settings for the BASH shell are kept in a file called .bash_profile or .bashrc in your home directory.

2.2 The Command Line

2.2.1 Accessing the Shell

The command line interface, provided by a shell program, is generally accessible through a terminal emulator program. The name terminal derives from the simple computers (terminals) that were used in the past to access large shared computer systems. When you start a terminal emulator you are generally dropped into your home directory. On Mac OS X you can use the "Terminal" program, while on a Linux systems with KDE you use "Konsole", on Linux systems with Gnome "GNOME Terminal" and on Linux systems with Unity the program is also called "Terminal". This chapter will only explain the basics of the BASH shell as it is standard on many Linux systems and on all recent versions of Mac OS X.

2.2.2 Navigating the File System

Assuming you have access to the command line, you will be presented with a so called prompt. What the prompt looks like depends on the version of UNIX or Linux you are using and how it was configured. In my case the command prompt looks like[2]:

---

[1]The example shows some directories that are Mac OS X specific: Applications, Developer, Library, Network, System and Volumes.

[2]What your prompt looks like exactly depends on the settings of your shell and is a matter of taste. My computer shows computername:directory username$, with Gretchen being the name of my computer.

```
Terminal — bash — 80×24
Last login: Mon Aug 11 16:23:59 on ttys004
Gretchen:~ thijscoenen$ cd coolstuff/
Gretchen:coolstuff thijscoenen$ ls
alice.txt        hello           noperm
Gretchen:coolstuff thijscoenen$ ▊
```

Figure 2.1: A screenshot showing Mac OS X's built-in Terminal program. This program allows you to access the underlying BASH shell on Mac OS X.

```
Gretchen:~ thijscoenen$
```

You can issue your commands by typing them followed by a press of the return key. One of the first things you may ask yourself, is where am I? Usually when you start a new command line session you will be dropped off in your user's home directory. To find out where in the file system you are you use the pwd (print working directory) command. The following example shows you what the output looks like for me:

```
Gretchen:~ thijscoenen$ pwd
/Users/thijscoenen
Gretchen:~ thijscoenen$
```

The home directory in these examples has a subdirectory called coolstuff that we will access. To change directories to the coolstuff directory use the cd (change directory) command. In the following example a pwd command issued as well to show we have actually changed directory:

```
Gretchen:~ thijscoenen$ cd coolstuff
Gretchen:coolstuff thijscoenen$ pwd
/Users/thijscoenen/coolstuff
Gretchen:coolstuff thijscoenen$
```

The eagle eyed among you may have noticed the ~ (tilde) appearing in the shell session examples. This is shorthand for your home directory, and your shell will expand it to your full home directory. You can use the tilde in commands to shorten them. To return to your home directory use the cd ~ command.

What does the coolstuff directory actually contain? To get a listing of the contents of a directory you can use the ls command. This command takes many options but the easiest way of using it is as follows:

```
Gretchen:coolstuff thijscoenen$ ls
alice.txt        hello               noperm
Gretchen:coolstuff thijscoenen$
```

Hidden files on UNIX systems have names that start with a dot. There are two special directories that are always present in every directory, the single dot and double dot directories. The former . is shorthand for the current directory and the latter .. is shorthand for the parent directory. By passing the -a option to ls you make hidden files visible.

```
Gretchen:coolstuff thijscoenen$ ls -a
.                 ..                alice.txt      hello               noperm
Gretchen:coolstuff thijscoenen$ cd ..
Gretchen:~ thijscoenen$ pwd
/Users/thijscoenen
Gretchen:~ thijscoenen$
```

The ls command has many options, but several are worthy of mention. First, the -l option will list the directory contents with one file or directory per line and also show the size (in bytes) and permissions for each entry. Second, the -h option will show the file sizes in human readable format (so not in bytes for large files, see below for an example).

```
Gretchen:coolstuff thijscoenen$ ls -l
total 312
-rw-r--r--  1 thijscoenen  staff  147731 22 jul 23:27 alice.txt
-rwxr--r--  1 thijscoenen  staff      43  7 aug 18:02 hello
-rw-r--r--  1 thijscoenen  staff      50 11 aug 14:35 noperm
Gretchen:coolstuff thijscoenen$ ls -lh
total 312
-rw-r--r--  1 thijscoenen  staff   144K 22 jul 23:27 alice.txt
-rwxr--r--  1 thijscoenen  staff    43B  7 aug 18:02 hello
-rw-r--r--  1 thijscoenen  staff    50B 11 aug 14:35 noperm
Gretchen:coolstuff thijscoenen$
```

Another nice command is ls -lrt which will list all files and sorted by the time of last access (oldest to newest). This is handy when you want to find the most recently accessed file in a directory as it will be at the end of the ls output.

The ls command can furthermore be used to list the content of a specific directory. Using the previous example I can list the contents of the coolstuff directory from my home directory using the command ls coolstuff. When you want to recursively list the contents of some directory you can use the ls -R command. It shows the content of the current directory, the content of the subdirectories and the content of the subdirectories of each sub directory etc. Because this command produces a lot of output I suggest you try it for yourself.

2.2.3 Running Programs

The shell can only start programs it can find, and the directories where the shell will look for programs are listed in an environment variable (see Section 2.2.13 for an explanation) called $PATH. For programs that are on the $PATH, i.e. in one of the directories listed by the $PATH variable, you can just type their name and press enter (even if the program is in a different directory than you are). The programs ls and cd are on the $PATH (on my system they actually reside in the /bin directory). When you try to run a non-existent program or a program that your shell cannot locate, you will see an error like the following:

```
Gretchen:coolstuff thijscoenen$ ls
alice.txt        hello              noperm
Gretchen:coolstuff thijscoenen$ nosuchthing
-bash: nosuchthing: command not found
Gretchen:coolstuff thijscoenen$
```

This will even happen for programs present in the same directory as you — unless that directory happens to be on the $PATH. To remedy this problem, prepend the program name with ./ (this tells your shell to look in the current directory).

```
Gretchen:coolstuff thijscoenen$ hello
-bash: hello: command not found
Gretchen:coolstuff thijscoenen$ ./hello
Hello world!
Gretchen:coolstuff thijscoenen$
```

You may not have sufficient permissions to run just any program on UNIX. To check whether you have permissions to run the program you can just attempt to run it, your shell will report an error if you have insufficient permissions.

```
Gretchen:coolstuff thijscoenen$ ./noperm
-bash: ./noperm: Permission denied
Gretchen:coolstuff thijscoenen$
```

All the previous examples were for programs that run for a short time, if you have long running programs you may want to continue issuing new commands without waiting for some program to finish. You could just start a second (or third etc.) terminal and continue issuing commands in that new terminal. The BASH shell however allows you to start programs in the background so that it is possible to keep issuing commands in that same shell. To start a program in the background just add a space and ampersand `&` to the command. You can move a program from the background to the foreground with the `fg` command:

```
Gretchen:background-demonstration thijscoenen$ ls
runs-1-minute
Gretchen:background-demonstration thijscoenen$ ./runs-1-minute &
[1] 55868
Gretchen:background-demonstration thijscoenen$ pwd
/Users/thijscoenen/background-demonstration
Gretchen:background-demonstration thijscoenen$ fg
./runs-1-minute
One minute passed
Gretchen:background-demonstration thijscoenen$
```

As you can see in the example above a number is shown in the shell after the program was started using the `&`. This number is the process number of the program just started, in Section 2.2.7 these process numbers will be explained. If you started a program in the foreground, but want to move it to the background you first suspend the program and then move it to the background. Suspend the program by pressing `Control` + `z` and then move it to the background by issuing the `bg` command.

```
Gretchen:background-demonstration thijscoenen$ ./runs-1-minute
^Z
[1]+  Stopped                 ./runs-1-minute
Gretchen:background-demonstration thijscoenen$ bg
[1]+ ./runs-1-minute &
Gretchen:background-demonstration thijscoenen$ pwd
/Users/thijscoenen/background-demonstration
Gretchen:background-demonstration thijscoenen$ fg
./runs-1-minute
One minute passed
Gretchen:background-demonstration thijscoenen$
```

In this example you can see that the `Control` + `z` key presses are shown in the shell as `^Z`. It is also possible to terminate a program running in the shell by pressing `Control` + `c`. Some programs may not react to an attempt to terminate it this way and you may have to resort to the `kill` command described below in Section 2.2.7.

Figure 2.2: The man documentation viewer showing part of the documentation for ls.

2.2.4 Getting Help

While this reader will get you started using Linux it cannot possibly explain all features of BASH or all programs available by default on a UNIX machine. Fortunately a large amount of documentation is available, some of it accessible directly from the command line and some of it on line. Most command line programs have simple built in help that is usually displayed when no options are specified on the command line or through options as -?, -h or --help.

While many programs provide some simple documentation about themselves UNIX provides a standard command line documentation viewer called man. To get help about a specific program just type the command man program-name, where you should replace program-name with the program that you want help for. Figure 2.2 shows the help for ls (accessed by typing man ls). The documentation can be navigated using the up and down arrows, to move up and down by a line, and space to jump a full page. The documentation can also be searched by typing a forward slash / followed by the word you are looking for and then return. When there are several matches for a search you can jump to the next match by pressing n (or in reverse with N). The viewer can be quit by typing a q. More information about navigating in man can be accessed by pressing h while it is running. The documentation available through man are also called "man pages". Some software will be documented using info, try running that like you would man in case no man pages are available (because the documentation may be written in info format).

On line there are several good resources for help with UNIX problems, that can be very helpful because the man pages are quite technical and at times hard to read. The very basic UNIX commands tend to have Wikipedia pages with examples. Stack overflow, a community site about programming problems (see http://stackoverflow.com), can be helpful as can be the related Stack Exchange site about UNIX and Linux (see http://unix.stackexchange.com). A nice site that can explain some shell commands is Explain Shell (see http://explainshell.com), it allows you to cut and paste a command and it then shows you an explaination of it. The Linux Documentation Project (see http://www.tldp.org) is also useful. A note about using Google to find documentation on the Internet: not everything you find will be correct so refrain from just copying whatever you find!

2.2.5   Copying, Moving and Removing Files

To copy files, or directories use the cp command. The general shape of the command is cp source destination. If the source is a file and the destination is also a filename the contents of source will be copied over the destination file. If destination does not exist yet source will be copied to a file with the specified destination as name. If destination is a directory the source file will be copied to it. If the source is a directory itself the copying will fail (these examples use cat to show file contents):

```
Gretchen:cp-demonstration thijscoenen$ ls -l
total 16
drwxr-xr-x  2 thijscoenen  staff  68 21 aug 22:40 directory_a
-rw-r--r--  1 thijscoenen  staff  16 21 aug 22:40 file_a.txt
-rw-r--r--  1 thijscoenen  staff  16 21 aug 22:40 file_b.txt
Gretchen:cp-demonstration thijscoenen$ cat file_a.txt
This is file A.
Gretchen:cp-demonstration thijscoenen$ cat file_b.txt
This is file B.
Gretchen:cp-demonstration thijscoenen$ cp file_a.txt file_c.txt
Gretchen:cp-demonstration thijscoenen$ ls -l
total 24
drwxr-xr-x  2 thijscoenen  staff  68 21 aug 22:40 directory_a
-rw-r--r--  1 thijscoenen  staff  16 21 aug 22:40 file_a.txt
-rw-r--r--  1 thijscoenen  staff  16 21 aug 22:40 file_b.txt
-rw-r--r--  1 thijscoenen  staff  16 21 aug 22:41 file_c.txt
Gretchen:cp-demonstration thijscoenen$ cat file_c.txt
This is file A.
Gretchen:cp-demonstration thijscoenen$ cp file_a.txt file_b.txt
Gretchen:cp-demonstration thijscoenen$ cat file_b.txt
This is file A.
Gretchen:cp-demonstration thijscoenen$ cp file_a.txt directory_a/
Gretchen:cp-demonstration thijscoenen$ ls directory_a/
file_a.txt
Gretchen:cp-demonstration thijscoenen$ cp directory_a whatever
cp: directory_a is a directory (not copied).
Gretchen:cp-demonstration thijscoenen$
```

To copy a directory and all it contains use the -r option of cp:

```
Gretchen:cp-demonstration thijscoenen$ cp -r directory_a/ whatever
Gretchen:cp-demonstration thijscoenen$ ls whatever/
file_a.txt
Gretchen:cp-demonstration thijscoenen$
```

To just move a file to a different filename or location use the mv command. Simple use of mv takes the shape mv source destination. If source is a file and destination does not exist yet source will be moved there (after this operation source will cease to exist). If destination does exist, and is a file, it will be overwritten. If destination is a directory the file source will be moved to that directory. If source itself is a directory it will be moved to destination, with similar rules:

```
Gretchen:cp-demonstration thijscoenen$ ls -l
total 24
drwxr-xr-x  3 thijscoenen  staff  102 21 aug 22:42 directory_a
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:40 file_a.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:42 file_b.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:41 file_c.txt
```

```
drwxr-xr-x  3 thijscoenen  staff  102 21 aug 22:45 whatever
Gretchen:cp-demonstration thijscoenen$ mv file_a.txt file_b.txt
Gretchen:cp-demonstration thijscoenen$ ls -l
total 16
drwxr-xr-x  3 thijscoenen  staff  102 21 aug 22:42 directory_a
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:40 file_b.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:41 file_c.txt
drwxr-xr-x  3 thijscoenen  staff  102 21 aug 22:45 whatever
Gretchen:cp-demonstration thijscoenen$ mv whatever directory_a/
Gretchen:cp-demonstration thijscoenen$ ls -l
total 16
drwxr-xr-x  4 thijscoenen  staff  136 21 aug 22:57 directory_a
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:40 file_b.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:41 file_c.txt
Gretchen:cp-demonstration thijscoenen$ mv directory_a/ file_b.txt
mv: rename directory_a/ to file_b.txt: Not a directory
Gretchen:cp-demonstration thijscoenen$
```

As you can see trying to move a directory to an existing file will fail.

Removing files is done with the `rm` command. Generally you just use the `rm somefile` command. If `somefile` is a directory this will fail because only *empty* directories are removed with `rmdir`.

```
Gretchen:cp-demonstration thijscoenen$ ls -l
total 16
drwxr-xr-x  4 thijscoenen  staff  136 21 aug 22:57 directory_a
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:40 file_b.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:41 file_c.txt
Gretchen:cp-demonstration thijscoenen$ rm file_b.txt
Gretchen:cp-demonstration thijscoenen$ ls -l
total 8
drwxr-xr-x  4 thijscoenen  staff  136 21 aug 22:57 directory_a
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:41 file_c.txt
Gretchen:cp-demonstration thijscoenen$ rm directory_a/
rm: directory_a/: is a directory
Gretchen:cp-demonstration thijscoenen$ rmdir directory_a/
rmdir: directory_a/: Directory not empty
Gretchen:cp-demonstration thijscoenen$
```

If you want to remove a directory and all its content (whether files or sub directories) you can use the `rm -rf *`. Note though that this is a very dangerous command, if not the "most dangerous command ever"[3], as it will recursively (the `-r` option) remove everything (matched by the `*` "glob pattern") and skip any questions (the `-f` option). For more information on glob patterns see Section 2.2.9. Note that `rm -rf *` cannot be undone!

2.2.6   Creating Directories

Because in the previous section it was shown how to remove directories you may ask yourself how you create them in the first place. UNIX has the `mkdir` command for that. You can create a directory with the command `mkdir somedirectory` where some directory is the name of the directory to be created. It is also possible to create several directories in one go by listing their names after `mkdir` separated with spaces. If you have to create a directory inside of several that is possible using the command `mkdir -p` followed the full path you want to create.

```
Gretchen:mkdir-examples thijscoenen$ ls
Gretchen:mkdir-examples thijscoenen$ mkdir dir1 dir2 dir3
```

---

[3]According to one of my proofreading colleagues.

```
●●●                    Terminal — top — 80×24
Processes: 62 total, 2 running, 60 sleeping, 267 threads        17:29:29
Load Avg: 0.13, 0.20, 0.17  CPU usage: 4.24% user, 9.43% sys, 86.32% idle
SharedLibs: 5196K resident, 5648K data, 0B linkedit.
MemRegions: 13931 total, 632M resident, 13M private, 228M shared.
PhysMem: 685M wired, 947M active, 311M inactive, 1942M used, 105M free.
VM: 156G vsize, 1039M framework vsize, 17259915(0) pageins, 6093270(0) pageouts.
Networks: packets: 64936199/83G in, 37656757/3714M out.
Disks: 28389042/396G read, 17131423/460G written.

PID    COMMAND     %CPU TIME     #TH #WQ #PORT #MREG RPRVT  RSHRD  RSIZE
58153  screencaptur 0.4 00:00.03 3   2   42    81    232K-  9000K+ 2812K
58151  top         5.8  00:00.33 1/1 0   25    33    984K+  264K   1564K+
58148  bash        0.0  00:00.00 1   0   17    24    392K   744K   1064K
58147  login       0.0  00:00.01 1   0   22    53    496K   312K   1596K
57871  vim         0.0  00:25.76 1   0   17    36    1692K  244K   3028K
57863  bash        0.0  00:00.01 1   0   17    24    408K   744K   1108K
57862  login       0.0  00:00.01 1   0   22    53    488K   312K   1596K
57789  bash        0.0  00:00.03 1   0   17    24    416K   744K   1116K
57788  login       0.0  00:00.01 1   0   22    53    488K   312K   1596K
57747  bash        0.0  00:00.04 1   0   17    24    424K   744K   1116K
57746  login       0.0  00:00.01 1   0   22    53    488K   312K   1596K
57729  bash        0.0  00:00.01 1   0   17    24    396K   744K   1064K
57728  login       0.0  00:00.01 1   0   22    53    488K   312K   1596K
57706  Preview     0.0  00:13.14 2   1   119   277   16M+   27M-   35M
```
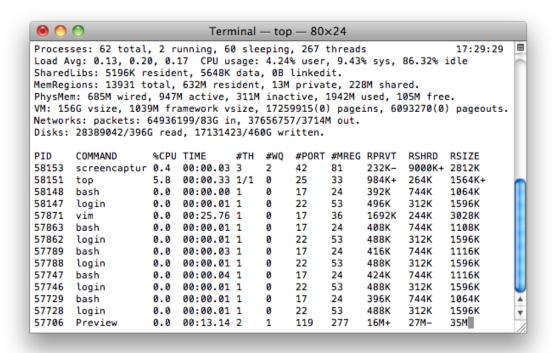
Figure 2.3: This screenshot shows top running on a Mac OS X system.

```
Gretchen:mkdir-examples thijscoenen$ ls
dir1        dir2        dir3
Gretchen:mkdir-examples thijscoenen$ ls -l
total 0
drwxr-xr-x  2 thijscoenen  staff  68 21 aug 23:28 dir1
drwxr-xr-x  2 thijscoenen  staff  68 21 aug 23:28 dir2
drwxr-xr-x  2 thijscoenen  staff  68 21 aug 23:28 dir3
Gretchen:mkdir-examples thijscoenen$ mkdir -p container/subdir/subdir/whatever
Gretchen:mkdir-examples thijscoenen$ cd container/subdir/subdir/whatever
Gretchen:whatever thijscoenen$ pwd
/Users/thijscoenen/mkdir-examples/container/subdir/subdir/whatever
Gretchen:whatever thijscoenen$
```

2.2.7   Listing and Killing Processes

On UNIX each running program consists of one or sometimes more processes that are identified by *process numbers* often abbreviated to PID. To get an idea of the running programs use the top command. This will show you a continually updating list of running processes, sortable by for instance memory use or processor use. When your computer is stuck it is quite often possible to find the offending program by looking through the most active processes in top. As with man you can exit top by pressing q. Figure 2.3 shows a screenshot of top. You can get a list of the currently running processes using the ps command, which unlike top will not continually update - it is like an ls for processes. Run without options ps only shows the processes running in the current shell while you can get a list of all processes running in all shells for a certain user using ps -u username (with username replaced with relevant username). For example, while logged in on a LOFAR[4] compute node as coenen:

---

[4] The Low Frequency Array, a large radio telescope array operating at low radio frequencies that has its core in the Netherlands.

```
coenen@locus048:~$ ps
  PID TTY          TIME CMD
14099 pts/2    00:00:00 bash
15518 pts/2    00:00:00 ps
coenen@locus048:~$ ps -u coenen
  PID TTY          TIME CMD
14098 ?        00:00:00 sshd
14099 pts/2    00:00:00 bash
15229 ?        00:00:00 sshd
15230 pts/3    00:00:00 bash
15266 pts/3    00:00:00 top
15522 pts/2    00:00:00 ps
coenen@locus048:~$
```

Another useful option to `ps` is `-A` which will select all processes, also not necessarily running in a terminal. The process numbers can be used to terminate unresponsive programs or those that are otherwise stuck. The `kill` command allows you to terminate processes by process number:

```
Gretchen:background-demonstration thijscoenen$ ./runs-1-minute &
[1] 59380
Gretchen:background-demonstration thijscoenen$ kill 59380
Gretchen:background-demonstration thijscoenen$ fg
-bash: fg: job has terminated
[1]+  Terminated              ./runs-1-minute
Gretchen:background-demonstration thijscoenen$
```

Some processes may still not terminate, you can then send them a stronger message using the `-9` (KILL) signal. E.g. to kill process `101` using the KILL signal: `kill -9 101`.

2.2.8 Dealing with Text Files

You will inevitably have to work with text files, because, among other reasons, programming source code are text files, many of UNIX's settings are in small text files in your home directory and quite a few scientific data sets are encoded in text files. Luckily UNIX has many tools to work with text files from the command line. The first command is `cat` which will show the contents of a text file in your terminal. Since `cat` does not paginate, this will likely cause your terminal to scroll. In the following example I look at the contents of the `hello` file (which is a Python script and therefore text):

```
Gretchen:coolstuff thijscoenen$ ls
Whatever        alice.txt       hello               noperm
Gretchen:coolstuff thijscoenen$ cat hello
#!/usr/bin/env python
print "Hello world!"
Gretchen:coolstuff thijscoenen$
```

If the file you are working with is large, then using `cat` to look at its content is inconvenient. By using the `more` or `less` commands you will be presented with a so-called pager that allows you to view and move around in the text file without scrolling. The `less` command allows you to move in both directions through the file using the up and down arrow keys. You can search for words by pressing a forward slash ⫶/⫶ followed by the word you are looking for and then ⏎enter⏎. If there are several matches you can move between them using ⫶n⫶ (or ⫶N⫶ for the reverse direction). Help is accessible by pressing ⫶h⫶ and you can quit `less` by pressing ⫶q⫶. This is very similar to the `man` works because it actually uses `less` to show the actual man page (visible in the title bar of the window in Figure 2.2).

If you are only interested in a quick peek at the contents of some text file, you can use the `head` or `tail` commands to, respectively, look at the first few lines or last few lines of a text file. Similarly if you want to find out the number of words or lines in a file the `wc` (word count) command is useful, without options it shows the number of lines, the number of words and the number of bytes in the file.
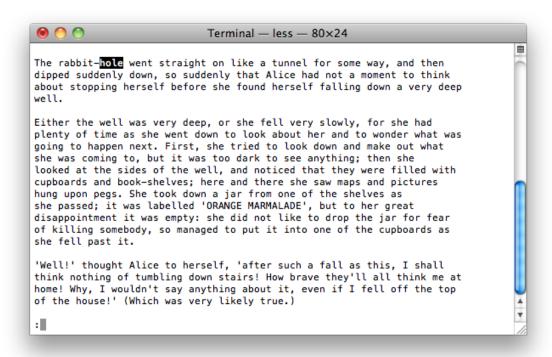
Figure 2.4: The less program showing a text file, you can move through the file using the arrow keys.

```
Gretchen:coolstuff thijscoenen$ head alice.txt
ALICE'S ADVENTURES IN WONDERLAND

Lewis Carroll


CHAPTER I. Down the Rabbit-Hole

Alice was beginning to get very tired of sitting by her sister on the
bank, and of having nothing to do: once or twice she had peeped into the
book her sister was reading, but it had no pictures or conversations in
Gretchen:coolstuff thijscoenen$

Gretchen:coolstuff thijscoenen$ wc alice.txt
    3335   26444  147731 alice.txt
Gretchen:coolstuff thijscoenen$ wc -l alice.txt
    3335 alice.txt
Gretchen:coolstuff thijscoenen$
```

The last example (the wc -l alice.txt command) shows you how to count only the number of lines in a file. A oft used file format to encode tabular data is a so called Comma Separated Values (CSV) file (white space or tab separation are also used frequently). The lines in these files generally correspond to one data point, you can thus get a quick feel for the size of the data set by using wc -l on that file.

So far we have only treated looking through files, not editing or creating them. Many different text editors are available on UNIX, spanning the gamut from editors that run in your terminal and use arcane commands to more modern text editors with a graphical user interface. Note though that in this context modern does not always mean more efficient! The choice of editor is personal and there is not enough space in this reader to explain how to use

them all. Instead, I have added a chapter to the appendix that lists some good text editors (see Appendix 20). For now I will just mention that you will want a text editor that is focused on programmers as they allow you to work more efficiently.

To search a file for a certain piece of text you can use the grep command. To find out on which lines the Cheshire cat appears in the text we use grep with the -n option to print line numbers and -i option to perform case insensitive matches:

```
Gretchen:coolstuff thijscoenen$ grep -ni "Cheshire" alice.txt
1313:'It's a Cheshire cat,' said the Duchess, 'and that's why. Pig!'
1319:'I didn't know that Cheshire cats always grinned; in fact, I didn't know
1433:was a little startled by seeing the Cheshire Cat sitting on a bough of a
1440:'Cheshire Puss,' she began, rather timidly, as she did not at all know
2068:'It's the Cheshire Cat: now I shall have somebody to talk to.'
2100:'It's a friend of mine--a Cheshire Cat,' said Alice: 'allow me to
2144:When she got back to the Cheshire Cat, she was surprised to find quite a
```

If you want more context to each match use the -B option to specify how many lines *before* a match should also be printed and -A for the number of lines *after* a match (output truncated after the first two matches to save space):

```
Gretchen:coolstuff thijscoenen$ grep -B 1 -A 2 -i "Cheshire" alice.txt

'It's a Cheshire cat,' said the Duchess, 'and that's why. Pig!'

She said the last word with such sudden violence that Alice quite
--

'I didn't know that Cheshire cats always grinned; in fact, I didn't know
that cats COULD grin.'


--
```

Each match is separated by a line with two dashes --. If you are only interested in knowing whether a file contains a match you can use the -l option to print the filename in stead of the matching text:

```
Gretchen:coolstuff thijscoenen$ grep -li "Cheshire" *.txt
alice.txt
Gretchen:coolstuff thijscoenen$
```

grep is a very powerful tool that can also match patterns, it is useful to read through the man pages to see what more it can do for you.

2.2.9 Glob Patterns

The BASH shell supports filename expansion, that is it will expand special wildcard symbols and match them to filenames. This is also called *globbing*, and the pattern that is matched *glob pattern*. I will first explain the parts that make up a glob pattern, then show you how you use them with BASH.

* Matches any text in filename.

? Matches zero or one unspecified character.

[ABC] Allow one character from a specified set of characters to be matched, in this case the characters A, B and C make up that set. Will match a single character only, not ABC.

[0-9] Will match a number in the range 0 through 9.

[a-z] Match any lower case character in the range a through z.

[A-Z] Match any upper case character in the range A through Z.

[A-Za-z] Match any upper case character in the range A through Z or any lower case character in the range a through z.

    \ Escape character, use this if you should match one of the characters that have a special meaning in a glob pattern. E.g. if you want to match the *, your glob pattern should include * so that the * itself matched and not expanded.

    ! This will negate a pattern. E.g. if you want to match anything not a lower case character use ![a-z].

Glob pattern can then be used in combination with other programs, what follows is a simple example that shows how ls can be combined with some glob patterns to look for some specific files.

```
Gretchen:glob-demonstration thijscoenen$ ls
101.txt                 604.txt                 ABC.txt
123.txt                 709.dat                 ABC123.txt
Gretchen:glob-demonstration thijscoenen$ ls [0-9]*
101.txt         123.txt         604.txt         709.dat
Gretchen:glob-demonstration thijscoenen$ ls [!0-9]*
ABC.txt                 ABC123.txt
Gretchen:glob-demonstration thijscoenen$ ls [5-9]*
604.txt         709.dat
Gretchen:glob-demonstration thijscoenen$ ls *dat
709.dat
Gretchen:glob-demonstration thijscoenen$ ls 1[2468]*
123.txt
Gretchen:glob-demonstration thijscoenen$
```

2.2.10   Finding Files

A common problem is, how do I find some specific file? The naive approach, use ls and move around the file system until you find what you were looking for, is time consuming and . Beyond the glob patterns that your shell provides (see Section 2.2.9) the UNIX find command is useful. With find you can look through a directory hierarchy for filenames matching a certain pattern or certain properties. You can then run some commands for each file that was found. At its simplest you can use find to look for a certain file called needle.txt in this example:

```
Gretchen:find thijscoenen$ find . -name "needle.txt"
./haystack/hay/hay/hay/needle.txt
./haystack/straw/needle.txt
Gretchen:find thijscoenen$
```

As you can see two files were found, but what if we know that the file we need contains the word gold? It turns out you can run other programs on each found file using the -exec option of find. In the following example grep is used to check whether a match contains gold:

```
Gretchen:find thijscoenen$ find . -name "needle.txt" -exec grep -li "gold" {} \;
./haystack/straw/needle.txt
Gretchen:find thijscoenen$
```

It is now clear that the file we needed was ./haystack/straw/needle.txt. In the commands that follow the -exec each occurrence of {} is replaced with the matching file name and each command must be followed by \;. The grep command was explained in Section 2.2.8.

2.2.11 `gzip` and `tar`

To conserve disk space it is sometimes a good idea to pack files more densely, there are several tools to do so. If you have used Windows so called zip files may be familiar. Although utilities to handle zip files exist (zip and unzip) on UNIX, you are more likely to come across `.gz` or `.bz2` files. Below I give examples of how to deflate and reinflate a file using gzip, note that I use a so-called pipe | and grep to only show filenames containing alice. Pipes are a way of sending output from one program to another, which are explained in Section 2.2.14. The following example shows that the Alice in Wonderland story can be packed to only about 36 % of its original size:

```
Gretchen:coolstuff thijscoenen$ ls -l | grep alice
-rw-r--r--  1 thijscoenen  staff  147731 22 jul 23:27 alice.txt
Gretchen:coolstuff thijscoenen$ gzip alice.txt
Gretchen:coolstuff thijscoenen$ ls -l | grep alice
-rw-r--r--  1 thijscoenen  staff  53596 22 jul 23:27 alice.txt.gz
Gretchen:coolstuff thijscoenen$ gunzip alice.txt.gz
Gretchen:coolstuff thijscoenen$ ls -l | grep alice
-rw-r--r--  1 thijscoenen  staff  147731 22 jul 23:27 alice.txt
Gretchen:coolstuff thijscoenen$
```

The bzip2 and bunzip2 programs work similarly and may achieve a slightly higher compression than gzip, but the latter is used more often.

When you download software or data you will come across so-called tarballs, which are files that can encapsulate many separate files or even directory full hierarchies. The name tarball derives from the tar utility that can create or unpack them. The name tar itself derives from tape archive (archiving to tape was and still is cheaper than archiving to disk). The example that follows shows how you create (-c option) a tarball that is compressed using gzip (-z option) how you can list its contents (-t option) and how you can extract it again (-x option):

```
Gretchen:coolstuff thijscoenen$ ls
Whatever        alice.txt       hello           noperm
Gretchen:coolstuff thijscoenen$ tar -czf cool.tar.gz *
Gretchen:coolstuff thijscoenen$ tar -tzf cool.tar.gz
Whatever
alice.txt
hello
noperm
Gretchen:coolstuff thijscoenen$ rm Whatever alice.txt noperm hello
Gretchen:coolstuff thijscoenen$ ls
cool.tar.gz
Gretchen:coolstuff thijscoenen$ tar -xzf cool.tar.gz
Gretchen:coolstuff thijscoenen$ ls
Whatever        cool.tar.gz     noperm
alice.txt       hello
Gretchen:coolstuff thijscoenen$ rm cool.tar.gz
Gretchen:coolstuff thijscoenen$
```

Note that some tarballs are compressed using bzip2 in which case you can use the `-j` option in stead of `-z`.

2.2.12 Permissions Continued

In Section 2.1.1 users and their permissions were introduced. UNIX has several commands that let you manage file permissions and ownership. UNIX also allows a user to temporarily acquire more privileges. We will start by demonstrating how file permissions can be changed. The owner of a file and a user with root privileges can change its permissions. The first step is to check the permissions, run for example `ls -l` to get a file listing that shows the permissions in the first column of the output.

```
Gretchen:chmod-demo thijscoenen$ ls -l
total 16
-rw-r--r--  1 thijscoenen  staff  68 27 aug 22:02 demo.py
----r--r--  1 thijscoenen  staff  15 27 aug 22:00 nopriv.txt
drwxr-xr-x  2 thijscoenen  staff  68 27 aug 22:12 subdir
```

The first column can be further split, taking demo.py as an example we get the file type -, owner permissions rw-, group permissions r-- and finally the permissions for everyone else r--. The file type can be: - for a file, d for a directory and l for a "symbolic link". In our example we are dealing with a normal file. The user permissions are listed in the order read, write and execute using the abbreviations rwx. If a permission is not granted the permission will be replaced with a dash. In our example the owner of the file has read and write permissions but no execute permissions: rw-. The group permissions in the example are only read permissions: r--. All other users are also only able to read the file: r--. The second column in the output of ls -l is not interesting for our current purposes, but the third is: it shows the owner of the file. In this case that user is thijscoenen. The following column shows the group that the file is assigned to, in this case the group staff. The next column shows the file size in bytes, followed by a column that shows the last time the file was changed. The last column shows the actual file or directory name. To change the permissions of the user, group and all other users the file owner (or a sufficiently privileged user like root) can use chmod. This command takes as options a string like ugo+rwx that specifies which user should gain or lose which permissions and finally as arguments the files for which the permissions should be changed. The options should be read as: the owner, the group and all others (ugo), gain (+) the permissions to read, write and execute (rwx). If we want to make the demo.py file executable for its owner, the following command will do the trick:

```
Gretchen:chmod-demo thijscoenen$ chmod u+x demo.py
Gretchen:chmod-demo thijscoenen$ ls -l
total 16
-rwxr--r--  1 thijscoenen  staff  82 27 aug 22:29 demo.py
----r--r--  1 thijscoenen  staff  15 27 aug 22:00 nopriv.txt
drwxr-xr-x  2 thijscoenen  staff  68 27 aug 22:12 subdir
Gretchen:chmod-demo thijscoenen$
```

If we want to disallow other users to see what is in the subdir directory we can remove (-) the execute permissions on that directory for its group and all other users:

```
Gretchen:chmod-demo thijscoenen$ chmod go-x subdir/
Gretchen:chmod-demo thijscoenen$ ls -l
total 16
-rwxr--r--  1 thijscoenen  staff  82 27 aug 22:29 demo.py
----r--r--  1 thijscoenen  staff  15 27 aug 22:00 nopriv.txt
drwxr--r--  2 thijscoenen  staff  68 27 aug 22:12 subdir
Gretchen:chmod-demo thijscoenen$
```

You can also specify several files by just typing their names or by using the appropriate glob patterns. To recursively change the permissions of some directory hierarchy use the -R option, for example to change the permissions on subdir and everything it contains to be completely open to everyone: chmod -R ugo=rwx subdir. In this last command we use = to just set the permissions. Old examples explaining chmod may show numerical permissions, but they are not necessary to manage the permissions and will not be treated here.

To change the owner of a file use the chown command. Using requires that you are either the owner or are otherwise sufficiently privileged. To change the owner of demo.py from the examples to root use the following command:

```
Gretchen:chmod-demo thijscoenen$ chown root:staff demo.py
chown: demo.py: Operation not permitted
Gretchen:chmod-demo thijscoenen$ sudo chown root:staff demo.py
Password:
Gretchen:chmod-demo thijscoenen$ ls -l
total 16
```

```
-rwxr--r--  1 root        staff  82 27 aug 22:29 demo.py
----r--r--  1 thijscoenen  staff  15 27 aug 22:00 nopriv.txt
drwxrwxrwx  2 thijscoenen  staff  68 27 aug 22:12 subdir
```

As you can see there is a small hitch, the user `thijscoenen` cannot change the ownership to `root` and group `staff`, because that user is not privileged enough. The second time `chmod` is preceded with `sudo`, which means execute the command that follows as a super user (`root` on most systems). To prove that you in fact are allowed to acquire `root` privileges you are asked to enter the password that belongs to that user. After acquiring the privileges `chown` to `root` becomes possible. While this example is a bit silly `sudo` is very important when installing software. A lot of software is installed system wide and that means somewhere in the system directories of UNIX. Those directories are not writable for normal user, a problem that is solved with `sudo`. Note: if you are using a university computer you will most likely not have the ability to use `sudo`, because most system administrators will not give out `root` privileges to just anyone — on your own Mac or Linux machine that should be no problem.

2.2.13   Environment Variables

Some of the settings of your BASH shell are put in so-called *environment variables*. A *variable* is a name (or an identifier) associated with some stored value that can be changed. We already came across an environment variable called `$PATH`, its value is a list of directories that are searched for executables. The BASH `echo $PATH` command will echo (show) the value of the variable `$PATH`, the `$` is needed to signify that the following name is a variable (the example below shows what goes wrong without the `$`):

```
Gretchen:~ thijscoenen$ echo $PATH
/Library/Frameworks/Python.framework/Versions/2.7/bin:/usr/bin:/bin:/usr/sbin:/sbin:/usr/local/bin:
/usr/texbin:/usr/X11/bin
Gretchen:~ thijscoenen$ echo PATH
PATH
Gretchen:~ thijscoenen$
```

As you can see `$PATH` contains multiple directories separated with a colon `:`. When you change `$PATH` you should know that the directories are searched for a executable program in the same order as they are specified in `$PATH`. To set a variable to a certain value use the `export` command, e.g. to set `$GREETING` to `Good morning` use the following command:

```
Gretchen:~ thijscoenen$ export GREETING="Good morning"
Gretchen:~ thijscoenen$ echo $GREETING
Good morning
Gretchen:~ thijscoenen$
```

This variable will be defined as long as your BASH session exists. If you are running several BASH sessions, the variable will only be defined for that one session where it was defined. To define a variable for all future BASH session you should edit the settings file for your BASH sessions. This file is probably called `.bashrc` (Linux) or `.bash_profile` (Mac OS X) in your home directory. Be careful editing these files, getting it wrong may make your shell inoperable.

If you want to know which environment variables are defined in your shell and what their values are you can use `env`. Running this command will result in a list of environment variables and their values. Because this is a long list I suggest you run this command on your own computer.

2.2.14   UNIX Pipes

So far we have treated a number of simple commands and utilities available on the command line. The power of the UNIX command line partially derives from the ability to connect different commands. The output of one program can be sent to another program using a so-called pipe. A pipe is signified with the `|` character. E.g. to use `grep` to look for a file called `.bash_profile` in my home directory:

```
Gretchen:~ thijscoenen$ ls -la | grep ".bash_profile"
-rw-r--r--   1 thijscoenen  staff      573 22 jul 23:38 .bash_profile
Gretchen:~ thijscoenen$
```

This may be a very simple example but the principle generalizes easily. The next example shows how you can search through a directory with some type of pulsar data files, looking for the dispersion measure, sorting by dispersion measure and then looking at only the first 10 lines:

```
Gretchen:00000143 thijscoenen$ ls *.inf | wc -l    313
Gretchen:00000143 thijscoenen$ cat *.inf | grep "Dispe" | sort -n -k 6 | head -n 10
 Dispersion measure (cm-3 pc)          =  2.6
 Dispersion measure (cm-3 pc)          =  2.65
 Dispersion measure (cm-3 pc)          =  2.7
 Dispersion measure (cm-3 pc)          =  2.75
 Dispersion measure (cm-3 pc)          =  2.8
 Dispersion measure (cm-3 pc)          =  2.85
 Dispersion measure (cm-3 pc)          =  2.9
 Dispersion measure (cm-3 pc)          =  2.95
 Dispersion measure (cm-3 pc)          =  3
 Dispersion measure (cm-3 pc)          =  3.05
Gretchen:00000143 thijscoenen$
```

Understanding how you can combine different UNIX command line utilities can help you automate boring tasks. This document is not about BASH programming so I will stop here now.

2.2.15   Output Redirects

The output of programs that shows up in your shell was written to "standard out" or "standard error" (a program should use the former for normal output while the latter can be used for errors or warnings). If you want these messages in a text file, you could copy the text from your shell into a text editor and save a file, but that is a lot of work. In stead you can just redirect standard out to a text file with a > sign:

```
Gretchen:coolstuff thijscoenen$ ls -l > blah.txt
Gretchen:coolstuff thijscoenen$ cat blah.txt
total 312
-rw-r--r-- 1 thijscoenen  staff        0 15 aug 16:35 Whatever
-rw-r--r-- 1 thijscoenen  staff   147731 22 jul 23:27 alice.txt
-rw-r--r-- 1 thijscoenen  staff        0 28 aug 01:31 blah.txt
-rwxr--r-- 1 thijscoenen  staff       43  7 aug 18:02 hello
-rw-r--r-- 1 thijscoenen  staff       50 11 aug 14:35 noperm
Gretchen:coolstuff thijscoenen$
```

If you want to capture both standard out and standard error in a file you can add 2>&1 to the end of the redirect command:

```
Gretchen:~ thijscoenen$ python divide_by_zero.py
Hi I'll divide by zero!
The result is
Traceback (most recent call last):
  File "divide_by_zero.py", line 2, in <module>
    print "The result is", 1/0
ZeroDivisionError: integer division or modulo by zero
Gretchen:~ thijscoenen$ python divide_by_zero.py > output1.txt
Traceback (most recent call last):
  File "divide_by_zero.py", line 2, in <module>
```

```
    print "The result is", 1/0
ZeroDivisionError: integer division or modulo by zero
Gretchen:~ thijscoenen$ cat output1.txt
Hi I'll divide by zero!
The result is
Gretchen:~ thijscoenen$ python divide_by_zero.py > output2.txt 2>&1
Gretchen:~ thijscoenen$ cat output2.txt
Hi I'll divide by zero!
The result is
Traceback (most recent call last):
  File "divide_by_zero.py", line 2, in <module>
    print "The result is", 1/0
ZeroDivisionError: integer division or modulo by zero
Gretchen:~ thijscoenen$
```

In this example you can see that Python sends the text associated with an `ZeroDivisionError` to standard error and the rest of the messages to standard out. The file `output1.txt` only contains the text that went to standard out, while the file `output2.txt` contains both the text for standard out and standard error. In case you do not want to overwrite an existing file, but append to the new output to it replace the > with >>.

Chapter 3

Python

The programming course of choice for this course is Python. It was originally developed by Guido van Rossum at the "Centrum voor Wiskunde en Informatica" (CWI) in Amsterdam and was made public in 1991. Nowadays Python is no longer developed at the CWI, but van Rossum is still leading its development. The Python syntax allows for concise code and the language is relatively easy to learn whilst still being powerful. Part of Python's popularity stems from its large standard library (Python is "batteries included") and the availability of a large number of external packages. For scientific programming in libraries like Numpy, Matplotlib and Scipy make Python a powerful tool.

There are two flavors of Python available at the moment, Python 2 and Python 3. Because most scientific software is still written in Python 2 and not all useful libraries are translated (*ported* in programming jargon) to Python 3. Eventually Python 3 will take over from Python 2, but that latter version will still be supported by the Python developers through at least 2020 (and probably longer by Linux distributors).

3.1   Programming Languages, Source Code, and Python

Computer programs can be created in a variety of different programming languages. The most hardware specific language is called *machine language*, it is the language that is directly executed by the computer hardware. To write programs in machine language a programmer would have to be intimately familiar with the details of the hardware, making the creation of even simple programs laborious. Fortunately programs can also be created in more abstract programming languages, that are then translated to machine language automatically using special computer programs. These special programs, compilers or interpreters (see below), know the details of the underlying hardware and how to interact with it.

The instructions in some programming language that specify a program's behavior, are called its *source code*. As mentioned above, this source code is later translated to machine code. Python is a *high level* programming language because it shields programmers from the underlying hardware almost completely, while other *lower level* programming languages may require knowledge of the hardware to a larger degree. The C programming language, for instance, requires one to do memory management (unlike Python).

The translation from source code to machine code can be performed in several ways. When the source code is translated to machine code in one go before the program is executed, the translation is called *compilation* and the program that does the translation a *compiler*. The aforementioned C language is generally compiled. When source code is translated instruction by instruction each time the program is run, it is said to be *interpreted* and the program that performs the translation an *interpreter*. The standard way of running Python programs uses an interpreter. The most well known Python interpreter, the one that is probably already installed on your computer if it runs Linux or Mac OS X, is referred to as CPython because that interpreter itself is written in C.

The short description of interpretation and compilation given above necessarily skips over many nuances and complications. Python code is actually compiled to *byte code* that is executed in a *virtual machine* — so-called because that program acts like virtual computer hardware. Several of the newer Python implementations mix interpretation with compilation in different ways that are (far!) outside the scope of this reader.

3.2   Running Python Programs

Python programs or scripts can be executed from the command line by starting the interpreter and passing it the name of the relevant Python file. We start with the classic "Hello world!" example without which no programming tutorial is complete. Start your text editor of choice, enter the following line of Python code and save it as `hello.py`.

```
print 'Hello world!'
```

This is about the simplest Python program you can produce that still does something. To run it enter the `python hello.py` command at your command prompt. Your shell session should look somewhat like mine:

```
Gretchen:python thijscoenen$ python helloworld.py
Hello world!
Gretchen:python thijscoenen$
```

As you can see all the program does is output the text `Hello world!` to your shell. The `print` statement sends the text that follows to the so-called "standard out", in this case your shell. Like all command line programs, Python programs can also accept command line options.
The Hello world example demonstrates one Python statement, `print`, that sends some text to the screen. In this case what it sends to screen is the *string* 'Hello world!'. A string is a basic data type of Python (representing a string of characters). The basic Python data types are treated in Chapter 4.

3.2.1   UNIX and the Shebang

It is possible on UNIX to avoid having to type `python` in front of the programs you want to run. To do so make the script executable (using the `chmod` command and add a so-called *shebang* to it. The shebang line informs UNIX which interpreter to use to run your program. The shebang must be the first line of a script and it looks as follows:

```
#!/usr/bin/env python
```

3.3   Interactive Python Sessions

Python can also be started without giving it a program to execute by entering the command `python` in your shell. Python will start a new session in interactive mode, you can recognize this mode by the prompt Python shows you: >>>. At this prompt you can enter some Python code to be executed. Recreating the previous example looks like:

```
Gretchen:python thijscoenen$ python
Python 2.7.3 (v2.7.3:70274d53c1dd, Apr  9 2012, 20:52:43)
[GCC 4.2.1 (Apple Inc. build 5666) (dot 3)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>> print 'Hello world!'
Hello world!
>>> exit()
Gretchen:python thijscoenen$
```

As you can see the output of the program is the same as before and it directly follows the line of code you wrote. In the following line the interactive interpreter is stopped using `exit()`. There is another way of stopping the interactive interpreter by pressing Control + d. The interactive interpreter gives immediate feedback, which makes it well suited for quick experimentation. As a side note: the interactive mode of Python is also referred to as the *Python REPL* for "read eval print loop" because it takes user input (the *read* step), evaluates it (the *eval* step) then prints the result (the *print* step) before "looping" back to the first step.
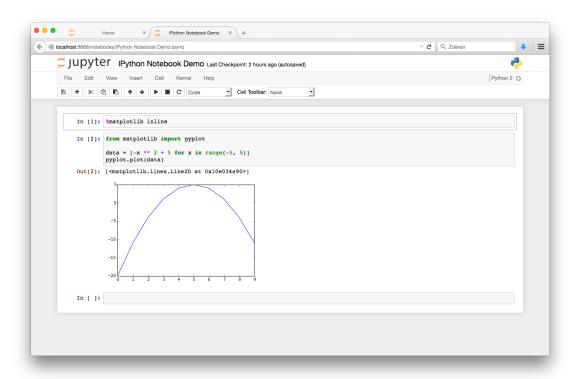
Figure 3.1: An example Ipython Notebook session demonstrating the ability to display graphs in notebooks.

3.4   Using Ipython Notebook

While Python comes with a built-in REPL, there are better ways of interacting with Python. IPython in particular is very convenient, it allows you to interact with Python and with the shell it is running under. IPython can be run in a terminal window, or you can run it as an IPython Notebook. In that latter mode IPython starts a local web server that you interact with through a web browser. IPython Notebooks are also able to display graphs using Matplotlib (see Chapter 9), mathematical equations (using LaTeX notation) and Markdown[1] formatted text.

When IPython Notebook is properly installed you can start it from the command line using `ipython notebook`. IPython Notebook will open a browser window and present you with a list of notebooks already open in the current directory (where you started IPython Notebook). Figure 3.1 shows a IPython Notebook session that displays a plot created using Matplotlib. The notebook viewer part of IPython Notebook is evolving into a more general tool named *Jupyter*. Jupyter can currently also be used with the R statistical programming language and the Julia programming language.

In an IPython notebook you can enter Python code and then executed it by pressing `shift` + `enter`. An IPython Notebook session behaves much like a "normal" interactive Python session but IPython adds a few magic functions. To invoke these IPython specific functions prepend the input line with a `%`. The previously mentioned example notebook shows the command that allows Matplotlib graphs to be shown (`%maplotlib inline`). Help about the available magic functions can is available by issuing the `?` command. The shell can be used directly from IPython Notebook by prepending a BASH command with an `!`. For example to use `ls` type `!ls`.

---

[1]`http://daringfireball.net/projects/markdown/`

Chapter 4

Variables and basic data types

The last chapter described several ways in which you can run snippets of Python, while this chapter describes what variables and data types are. It furthermore explains the basic data types of Python and give some examples of how they are used.

4.1 Variables and how to define them

Let's say you want to calculate the travel time of light from the Sun to Earth in minutes. As was shown in the last chapter you can start Python in its interactive mode and enter the calculation. Using 150 million kilometers for the distance of the Sun to Earth and 300 thousand kilometers per second as the speed of light. The division operator in Python is written as a single forward slash /. Remember to press ⌷Enter⌷ after entering a line of Python code.

```
>>> (150000000 / 300000) / 60.
8.333333333333334
```

Writing small snippets like this is fine for little experiments. You may, however, need the speed of light more often than just once, and repeatedly entering its value is tedious and error prone. It is better to enter the value once and associate a name with it, that you can then refer to when you need that value again. Variables in Python allow you to do just that, you choose a name and assign a value to it. Python's *assignment operator*, represented by a single equals sign =, is used to assign a value written to the right of the operator to a name written to the left of that operator. Python has no concept of constants and it is therefore always possible to change the value associated with a name by assigning it a new value (Python has no constants, unlike some other programming languages). Using variables the previous snippet can be rewritten using variables as follows:

```
>>> c = 300000
>>> d = 150000000
>>> seconds_per_minute = 60.
>>> (d / c) / seconds_per_minute
8.333333333333334
```

This example may be contrived but is shows you how variables are defined and then used in a calculation. Arithmetic will be described in some detail below.

4.2 Data types

One important property of values was skipped over in the previous section, namely their type. In Python every value has a type which describes what the value represents. Are we dealing with a number, a piece of text, a list of things, or something completely different? This is important because in memory these are all represented as a series of ones and zeros!

In Python, when you assign a new value to a variable the type of the value that the variable points to can change. Python is said to have dynamic types. Not all programming languages share this property. The C language, for

instance, does not allow you to change the type associated with a variable and is therefore said to have static types. In C a variable is of a certain type, in Python it is the values that have a type. In Python a variable is properly understood as a label pointing to some value, and the type of a variable is really is the type of the value pointed to. In Python it is always possible to find out what type is associated with a variable by using the type function. Functions will be explained in more details later, but for now you can see them as named pieces of a program that operate on a set of inputs to provide some return some outputs. The type function can be used, "invoked" or "called" in programming jargon, by writing its name and adding parentheses around some input value(s). The type function can be called as type(someval) to find the type associated with the someval variable. Some examples below (reusing the previous definitions for the c and seconds_per_minute variables):

```
>>> type(10)
<type 'int'>
>>> type(c)
<type 'int'>
>>> type(seconds_per_minute)
<type 'float'>
```

This example shows two types of variables, both representing numbers, the so-called *integer* numbers (called int for short) and floating point numbers (float for short). These are two of the basic data types available in Python that will be described in more detail below.

While Python is dynamically typed, it is strict in the way you operate on values of different types. Trying to, for instance, add a piece of text to a number using the addition operator + will result in a so-called TypeError. The following snippet shows what happens when you add a string (the data type that represents text) to an integer:

```
>>> '1' + 1
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
  TypeError: cannot concatenate 'str' and 'int' objects
```

Some programming languages, such as JavaScript or BASH, will perform an implicit type conversion to arrive at a result for this kind of operations — Python does not. While the Python behavior may seem strict, it is consistently applied and that makes programs easier to debug.

4.3   Mutable and immutable values in Python

When a value can be changed after it is created it is called *mutable*. When a value cannot be changed after it is created it is called *immutable*. Even "immutable variables" (in quotes because the values are immutable) can get a new value by assigning them a new value. What will not happen, however, is that the underlying value is changed. A simple example of immutable values is the Python string data type used to represent text. If you use the upper function associated with a string you will get an upper case version of that string, but the original string will not change, as is shown in the following snippet:

```
>>> s = 'Hello world!'
>>> s.upper()
'HELLO WORLD!'
>>> s
'Hello world!'
```

A similar simple example of a mutable data type is the Python list data type, which acts as a container for an arbitrary number of values. The sort function associated with lists will change the original list and, unlike the upper function shown above, it will not return a new copy of the list.

```
>>> l = [10, 5, 9, 0]
>>> l
[10, 5, 9, 0]
>>> l.sort()
```

```
>>> l
[0, 5, 9, 10]
```

Mutability versus immutability is frequent source of misunderstandings and errors in Python, so be careful.

4.4 Literals

This section introduces one last piece of jargon (for now) before continuing on with more practical examples. In the previous sections you saw how variables are defined and you were given a high level description of why types matter. What you also saw was how values are represented in snippets of Python source. The representation of a value in source code is called a *literal*. You saw integer literals, floating point literals, string literals and a list literal in the previous sections.

4.5 Numbers and basic arithmetic in Python

In Section 4.2 two of Python's data types, namely integers and floating point numbers, made a brief appearance. Integers represent the whole numbers including zero and negative numbers, while floating point numbers approximate the real numbers. I say "approximate" because in Python floating point numbers generally use 64 bits, allowing for only 2^{64} different numbers to be represented. Python's integers are special in that they can change internal representation in such a way that allow arbitrarily large numbers to be represented (or in the words of the Python documentation its integers numbers represent values "which are limited only by available memory"[1]).

Integers have a number of literal notations that you may come across. All previous examples showed integers literals written in their decimal form, there are also octal and hexadecimal forms. Octal and hexadecimal representations use respectively base 8 and base 16 to write the numbers. Octal digits run 0 through 7 while hexadecimal digits run 0 through 9 and then A through F. To differentiate these notations from decimal notatation octal numbers are preceded by 0 and hexadecimal numbers are preceded by 0x.

```
>>> # octal (note the leading 0):
... 011
9
>>> # hexadecimal (note the leading 0x):
... 0xff
255
```

Hexadecimal notation is quite often used in more low level (i.e. closer to the hardware or less abstract) code to represent numerical constants. You will certainly come across memory adresses written in hexadecimal notation in Python error messages. The # signs, also known as a hash or octothorpe, is used to signify a *comment*. Comments start after the # and extend to the end of the line, they are added to help programmers better understand the code and are ignored when the Python program is executed. The ... in this example are the way Python interactive sessions show that several lines were entered before they were executed and the result shown.

Like integers floating point numbers have several notations, they can be in their decimal form using a decimal dot to differentiate them from integers and they can be in scientific form. The following example shows these notations:

```
>>> # decimal notation:
... 3.1415
3.1415
>>> # scientific notation:
... 3e+5
300000.0
>>> 6.7e-11
6.7e-11
>>> 1e-2
0.01
```

---

[1] See https://www.python.org/dev/peps/pep-0237/.

```
>>> 1E-2
0.01
```

Python supports a small set of standard operators for numbers, while a larger set of mathematical functions are available through a special module called `math`. This section only treats the former, deferring the latter.

Addition in Python uses the + operator, subtraction uses the - operator, multiplication uses the * operator while raising to the power uses the ** operator. Standard precedence order is used in Python meaning that raising to the power is performed before a multiplication, which in turn is performed before a subtraction or addition. The full list of operator precedences is available from the Python website[2]. If in doubt, use parentheses to explicitly control the order in which the operators are applied yourself.

The division operator was deliberately not mentioned yet, because there are a few important details you should know about. The normal division operator in Python written as / has behavior that depends on the types that it operates on. Divide two floating point numbers and it will return a floating point number as answer, when it operates on a floating point number and an integer number it will return a floating point number, but when it operates on two integer numbers it will return an integer and throw away the remainder (the answer is *truncated*). This behavior is inherited from the C programming language and is a frequent source of bugs in Python programs. In Python 3 (*i.e. the version of Python that this reader is emphetically not about*) the division operator was split into two operators. The first division operator written / returns a floating point number regardless of whether it operates on floating point numbers or integers. The second division operator written // always truncates regardless of whether it operates on floating point numbers or integers. This behavior can be turned on also for Python 2 by adding the line

```
from __future__ import division
```

to the top of your programs (before all other import statements). Other than the divsion operator(s) of Python there is also the modulo operator written %. This operator returns the remainder of a division. Some examples of the default behavior of the Python division operator are shown below:

```
>>> 10 / 2
5
>>> 10 / 3
3
>>> 10 % 3
1
>>> from __future__ import division
>>> 10 / 2
5.0
>>> 10 // 2
5
>>> 10 / 3
3.3333333333333335
>>> 10 // 3
3
```

To solve truncation issues in programs that do not use the special `import` line you must manually convert the type of the operands. Python has built in functions to convert to and from strings, integers and floats. The `float` function converts its input to a floating point number and the `int` function converts to an integer. By converting one of the numbers in a division to a `float` the answer will not be truncated:

```
>>> float(3)
3.0
>>> 10 / float(3)
3.3333333333333335
>>> int(3.1415)
3
```

---

[2]See `https://docs.python.org/2/reference/expressions.html#operator-precedence`.

4.6 Strings in Python

Text of any length is represented as a string (type `str`). String literals can be delimited with single quotes (`'abc'`), but double quotes are also allowed (`"abc"`). A number of characters cannot be represented directly in a string literal and must be written as a so-called *escape sequence*. These sequences start with a \ and the character that follows is interpreted in a special way. An example is the tab character which may look like several spaces but is in fact only one character \t, the new line character is another example \n. A table of string escape sequences can be found in the section about string literals in the Python documentation[3]

Beyond the single and double quotes you can also come across triple quoted strings in Python source code (useful because they can contain new lines and quotes without escaping them) or strings broken up across several lines, see the following string definitions:

```python
# Several different string literals
s1 = 'abcdefg'
s2 = "abcdefg"
s3 = '''This string
contains 'quotes' and
new lines!
'''
s4 = """This is also allowed"""
```

The normal Python 2 strings use 8 bits to store each character and thus only 256 characters can be encoded. That is clearly not enough to handle all the world's characters or languages directly. Unicode was designed to solve this problem once and for all. It allows all the world's known characters to be represented and it can thus be used throughout the world. Python has a special Unicode string data type (`unicode`). Unicode literals are delimited with `u'abc'` and special characters can be entered using Unicode escapes that start with \u. The escape sequence for the Euro-sign, for example, is \u20ac.

```python
>>> u = u'abcd\u20ac'
>>> type(u)
<type 'unicode'>
>>>
```

4.6.1 Indexing and slicing

Each individual character in a string has an index that can be used to access it. These indices start at 0 and run to `n-1` for a string of `n` characters length. For example, you can access the first character of a string `s` as `s[0]` and its second character as `s[1]`. If you use an index that is too high, Python will raise an `IndexError`. As a convenience Python allows negative indices that are used as indexes starting at the end of the string. The last character of a string `s` can be accessed as `s[-1]` while the second-to-last can be accessed as `s[-2]`, etc.

To extract more than just one character from a string the slice operator is used. For a string `s` it is written as `somestring[startindex:endindex]`, where `startindex` and `endindex` are respectively the index of the letter where the slice starts and the index of the character before which the slice ends. To extract a sub string starting at the first character the `startindex` can be left out, similarly if the `endindex` is left out the sub string runs to the last character.

```python
>>> s = 'abcdefg'
>>> type(s)
<type 'str'>
>>> s[0]  # normal indexing
'a'
>>> s[10]  # accessing a non-existant character
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
IndexError: string index out of range
```

---

[3]See `https://docs.python.org/2/reference/lexical_analysis.html#strings`.

```
>>> s[-1]  # accessing the last character
'g'
>>> s[0:5]  # accessing the first 5 characters
'abcde'
>>> s[:5]  # idem
'abcde'
>>> s[5:]  # access starting at the 6th character
'fg'
```

Python implements many standard string manipulations as special functions connected to the string objects[4]. Below some examples of useful string manipulations; more are available in Python[5]. As was mentioned before, Python strings are immutable and can thus not be changed after they are created. These functions therefore create new strings in stead of changing the original strings they were called on.

```
>>> s.startswith('abc')
True
>>> s.endswith('xyz')
False
>>> s.upper()
'ABCD EFG'
>>> s.isdigit()
False
>>>
```

Because a sequence of characters are contained in strings and unicode strings, they are so-called *sequence types*. Indexing and slicing operations are also supported by other sequence data types in Python like tuples and lists (see below).

4.7   Tuples

A tuple is a container for several other values, whose types can be mixed in a tuple. The separate values contained in a tuple can be accessed by their indices. Because a Python tuple is immutable, it is not possible to assign new values to some position in a tuple. If you do try to assign new values to a tuple a TypeError will be raised. In source code tuple literals are usually enclosed in parentheses () and the values in the tuple separated by commas. Note that you can leave the parentheses out but not the comma(s). Like strings, tuples are a sequence type. The following example shows some operations on tuples.

```
>>> t1 = (1, 2, 3, 'abc')
>>> type(t1)
<type 'tuple'>
>>> t2 = 1,
>>> t2
(1,)
>>> type(t2)
<type 'tuple'>
>>> t3 = (1)
>>> t3
1
>>> type(t3)
<type 'int'>
>>> t1[0]
```

---

[4]I am avoiding the term *method* here, because that is associated with object oriented programming. But since Python is object oriented and its strings objects these functions are properly called methods.

[5]See the documentation at https://docs.python.org/2/library/stdtypes.html#string-methods

```
1
>>> t1[0:]
(1, 2, 3, 'abc')
>>> t1[0] = 0
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: 'tuple' object does not support item assignment
>>>
```

4.8 Lists

Like tuples and strings, lists are a sequence type that can contain any number of values of any type. Unlike tuples, lists can be changed after they are created (both their contents and their lengths can be changed), they are therefore *mutable*. Lists literals are enclosed in square brackets [] and the values in a list are separated by commas. Lists like strings support indexing and slicing to access values or extract parts of a list. Like strings, lists have many useful methods and because lists are mutable they are changed by these methods. The following example demonstrates some list methods:

```
>>> l = [1, 2, 3, 4]
>>> type(l)
<type 'list'>
>>> l.pop(0)
1
>>> l
[2, 3, 4]
>>> l.append(1)
>>> l
[2, 3, 4, 1]
>>> l.sort()
>>> l
[1, 2, 3, 4]
```

4.9 Dictionaries

Dictionaries are used to store key-value pairs, that is they map a certain key to the value that was stored with that key. Unlike lists that use indices to access stored values dictionaries use keys. Values can be any Python object but the keys have to be "hashable" (what that means exactly we will skip for now, but in practice integers, floats and strings can used as keys). Unlike tuples and lists, dictionaries are not ordered and you cannot speak of its first or last element in a dictionary. Trying to use a key that is not present in the dictionary will result in a KeyError being raised. Dictionaries correspond to *hash tables* in other programming languages.

```
>>> d = {3: True, 1: 'a', 5: 123}  # a dictionary literal
>>> type(d)
<type 'dict'>
>>> d[3]  # accessing an existing key
True
>>> d['no such key']  # accessing a non-existant key
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
KeyError: 'no such key'
>>> del d[1]  # removing a key-value pair (using the key)
>>> d
{3: True, 5: 123}
>>>
```

4.10   Booleans

Python has a special data type to represent truth values so called booleans (type `bool` for short). Booleans can only take one of two values, a boolean is either `True` or `False`. The results of comparisons are booleans (for more see Section 5.1):

```
>>> b = True
>>> type(b)
<type 'bool'>
>>> 1 > 2
False
>>> 2 > 1
True
>>>
```

4.11   NoneType

To signify a missing value Python uses the special `None` object of type `NoneType`. This is very useful to differentiate a missing value `None` versus for instance an empty string `""` or an empty list `[]`.

4.12   Basic type conversions

As was mentioned above Python is quite strict about its types and usually does not convert between different types automatically. Strings cannot be added to integers, mathematical functions will not work on strings, etc. When you use the wrong types typically a `TypeError` is raised. Fortunately you can convert between different types by using the right functions. To convert a string or a float to an integer use the `int()` function, to convert something to a float use `float()`, to convert to string use `str` and to convert to a Unicode string use `unicode()`. If such a conversion is not possible, a `ValueError` will be raised (for example *'xyz'* cannot be converted to a floating point number while *'1.5'* can.

```
>>> int(1.5)
1
>>> int('1.5')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: invalid literal for int() with base 10: '1.5'
>>> float('1')
1.0
>>> float(2)
2.0
>>> str(4.576)
'4.576'
>>>
```

For tuples, lists and dictionaries similar functions are available. The conversion functions for tuples and lists, respectively `tuple()` and `list()` expect a so-called *iterable* (we forego a exact definition at the moment but mention that an iterable contains several values that can be accessed one-by-one in a process called iteration). Dictionaries are a bit more complicated, as they need key-value pairs, so the `dict()` function expects an iterable of key-value pairs (so each object in the iteration needs to have two values). This "iterable of key-value pairs" is in practice often a list of tuples or a list of lists, where the inner tuples or lists have the key as the first element and the corresponding value as the second element.

```
>>> list('1234')
['1', '2', '3', '4']
```

```
>>> tuple('abcd')
('a', 'b', 'c', 'd')
>>> tuple([1, 2, 3, 4])
(1, 2, 3, 4)
>>> list((1, 2, 3, 4))
[1, 2, 3, 4]
>>> dict('1234')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: dictionary update sequence element #0 has length 1; 2 is required
>>> dict([(1, 'a'), (5, 'b')])
{1: 'a', 5: 'b'}
>>> list({1: 'a', 3: 'c'})
[1, 3]
>>>
```

The last example, converting a dictionary to a list, discards the values in the dictionary being converted because iteration over a dictionary returns only its keys. We will revisit iteration in Section 6.

Chapter 5

Conditional execution

When some piece of Python code is only executed when a certain condition is met that piece of code is *conditionally executed*, conditional execution is a fundamental concept in any programming language. This chapter explains how conditions are expressed in Python and used to control program flow.

5.1 Boolean expressions

An expression that evaluates to either `True` or `False` is a called *boolean expression*. Comparisons are common boolean examples of a boolean expressions

```
>>> a = 10
>>> b = 5
>>> a == b  # Check whether a equals b
False
>>> a != b  # Check whether a does not equal b
True
>>> a > b  # Check whether a is greater than b
True
>>> a >= b  # Check whether a is greater than or equal to b
True
>>> a < b
False
>>> a <= b
False
```

5.2 Boolean operators **and**, **or** and **not**

Boolean expressions can be combined using the and and or keywords and negated using the not keyword. When you have two booleans (or boolean expressions) the and operator will result in `True` if both booleans are `True` else the result will be `False`. The or operator results in `True` if either or both the booleans are `True` else the result will be `False`. The not operator negates a boolean turning `True` into `False` and vice versa.

```
>>> t = True
>>> f = False
>>> t and t
True
>>> t and f
False
>>> t or t
```

```
True
>>> t or f
True
>>> not t
False
>>> not f
True
```

When several boolean expressions are combined using the and, or or not the resulting expression is still a boolean expression because it evaluates to True or False.

5.3   Truthy and falsy

While Python has proper booleans True and False there are several other values that are treated as true or false in the context of a conditional. These are sometimes called *truthy* for values that are treated as True and *falsy* for values treated as False. Empty lists, dictionaries, sets and strings are falsy, whilst they are truthy when they have elements. In the following example the bool built-in function is used to test the truthiness of an empty list and a list with elements:

```
>>> l1 = []
>>> l2 = [1, 2, 3, 4]
>>> bool(l1)
False
>>> bool(l2)
True
```

5.4   `if`, `elif` and `else`

The if statement allows the conditional execution of a block of code. In Python blocks of code are defined by indenting them. The following trivial example shows that if followed by a True or False will either always be executed or never:

```
if True:
    print 'This will be printed'
if False:
    print 'This will never be printed'
```

Using a boolean expression in stead of True or False directly allows you to test a condition and only execute the following block if the boolean expression evaluates to True.

```
x = 10
y = 20

if x > y:
    print 'x is greater than y'
else:
    print 'x smaller than or equal to y'
```

This example also demonstrated the else clause, which will be executed only when the preceding condition (defined by the if statement) is False. When several conditions must be tested you could do that by repeatedly using if statements, but Python also has the elif statement. The elif statement can not occur by itself, it is always preceded by another elif statement or an if statement. Similarly else can only be used following an elif or if statement.

```
x = 10

if x < 0:
    print 'x is negative'
elif x == 0:
    print 'x is zero'
else:
    print 'x is positive'
```

There is a difference between using several `if` statements and using only one `if` and one or more `elif` statements. In the former case all conditions are checked while in the latter case the checking stops as soon as a condition is `True`.

Chapter 6

Loops and iteration

6.1 Introduction

When you need to repeat some instructions several times you can create a so-called loop to do this. There are two types of loops in Python, `while`-loops and `for`-loops. The former executes a block of code as long as a certain condition holds true, and the latter performs an action for each of the vales of an object like a list or dictionary. Each time the code block is executed is also called an *iteration*. The process of accessing the elements in a list is also called iteration, so the term can be a confusing. This chapter explains how the basic sequence types of Python can be iterated over.

6.1.1 The `while` loop

The `while` loop is executed as long as the condition at the beginning of the `while` loop is `True`. This condition is given as a boolean expression much like those used with the `if` and `elif` statements. The code that is to be repeatedly run is put in an indented block, again like the code blocks for `if` and `elif` statements. To break out of a loop you can use the `break` statement, when break is encountered the current loop is stopped immediately. If a loop is nested inside another loop breaking out of the inner loop will not break out the outer loop. The `continue` statement stops the current iteration immediately and jumps to the next iteration. A feature of loops in Python, one that is not shared with many other languages, is the possibility to add an `else` statement to a loop. The block of code associated with the `else` statement will be executed only if the `while` loop exits normally (that is without the use of `break`).

```python
i = 0

while i < 10:
    print i
    i = i + 1
else:
    print 'i is equal or larger than 10'

while True:
    pass  # Infinite loop, nevers stops (unles you kill the program) !!!
```

6.2 The `for` loop and iteration

The Python `for` loop always iterates over some object, i.e. it accesses all the items in that object one-by-one[1] and runs a block of code. When all elements have been accessed, or a `break` statement was encountered the iterations

---

[1]It is possible in Python to write functions or classes that can be iterated over without them having elements, much like the built-in `xrange` function. In this section we will not consider these types of iterables.

stops. The `continue` statement can be used with `for` loops as well, and functions in the same way as it does for `while` loops. Like `while` loops it is possible to use the `else` statement in combination with a `for` loop. The block of code associated with the `else` statement is executed when all items were accessed without encountering a `break` statement.

```python
s = 'abcdefg'

for character in s:
    if character == 'a':
        continue  # 'a' will never be printed because of the continue
    elif character == 'z':
        break  # if a 'z' is encountered the loop is exited
    else:
        print character
else:
    print 'There was no z in this string.'
```

Objects, like lists, that can be iterated over are called *iterable*. While in some languages `for` loops always use indices they should be avoided in Python (unless you really need the indices themselves). It is possible to iterate over the indices of for instance a list by using the `range()` function.

```python
s = 'abcdefg'

for i in range(len(s)):
    print s[i]
```

6.3   Iteration over basic Python types

6.3.1   Iterating over lists

```python
l = ['abc', 'def', 'efg']

# print elements of l one-by-one
for element in l:  # preferred!
    print element

for i in range(len(l)):
    print l[i]

# print index and corresponding element, element-by-element
for i, element in enumerate(l):  # preferred!
    print i, element

for i in range(len(l):
    print i, l[i]
```

6.3.2   Iterating over lists of tuples

```python
# Iteration over a list of tuples:
l1 = [('a', 'b'), ('x', 'y'), ('p', 'q')]

# print pairs of elements:
for c1, c2 in l1:  # preferred!
    print c1, c2
```

```python
for tup in l1:
    print tup[0], tup[1]

# This style of iteration also works with more than two element tuples
l2 = [(1, 2, 4), (2, 3, 5), (6, 2, 9)]

# print the elements of each 3-tuple:
for a, b, c in l2:
    print a, b, c
```

6.3.3   Iterating over strings

```python
s = 'abcdef'

# printing characters in s one-by-one
for character in s:  # preferred!
    print character


for i in range(len(s)):
    print s[i]


# printing index and corresponding character, character-by-character
for i, character in enumerate(s):  # preferred!
    print character, i


for i in range(len(s)):
    print i, s[i]
```

6.3.4   Iterating over dictionaries

```python
# dictionary iteration:
d = {1: 'a', 0: 'b', 6: 'c'}

# iteration over dictionary keys:
for key in d:
    print key


# iteration over dictionary values:
for value in d.values():
    print value


# iteration over key-value pairs:
for key, value in d.items():  # preferred (more readable)!
    print key, value


for key in d:
    print key, d[key]
```

Chapter 7

Simple File Input and Output

Python has a built-in `File` object that is used to represent an open file. To open a file you use the `open` function and to close it you use the `close` method of file objects (examples below). A file can be opened in one of several *modes*, it can be opened for *reading*, for *writing* or for *appending*. Reading leaves a file unchanged, while writing will change the file's content and appending will add to the end of a file.

7.1   Text files

The `open` function takes two arguments, first the name of the file as a string, and second the mode that the file should be opened in as a string. That latter string can be one of: `'r'` for reading, `'w'` for writing, `'a'` for appending. These modes work for textual data, binary data has special modes that will not be treated here — their documentation is available from the Python website. The `open` function returns a file object. Depending on the mode it is opened in this file object can be written to or read from using respectively the `write` or `read` methods of the file object.

Since operating systems have limits on the number of files that can be simultaneously opened, you must close files that you open. Python will automatically close open files when your program terminates, but that may be too late to avoid errors.

In the next example a small text file is created, then written to, and finally closed. The file object in this example must be assigned to a variable, `f` in this case, because you need a reference to it so that you can write to it and close it. The second part of the example shows how all data can be read from a text file.

```
# Open a file for writing and close it again:
f = open('smallfile.txt', 'w')
f.write('Some text for this file\n')
f.close()


# Now open the example file for reading, and print out the file content:
f = open('smallfile.txt', 'r')
print f.read()
f.close()
```

Note, this is the simplest example of writing and then reading a file, but it is not the recommended way of doing it — see below for a better technique.

Modern Python programs use the `open()` function in conjunction with the `with` keyword. The `with` keyword in the following example will ensure that the open file is closed no matter what, when the block defined by `with` is left. You will no longer need to close the file manually. A further nice property of Python `File` objects is that they can be iterated over (in the case of text files). Iterating over an open text file will access the lines in that file one-by-one. The following example demonstrates both the `with` statement and iteration over a text file.

```
with open('alice.txt', 'r') as f:
    count = 0
```

```
    for line in f:
        count = count + 1

print count
```

Chapter 8

Program Structure Basics

The previous chapters have shown you parts of Python, but not what a Python program actually looks like. This chapter will explain how you can organize a program in such a way that it is readable, reusable and well documented. A well structured program is a modular program with constituent parts that each have well defined purposes and is as independent from the rest of the program as possible. By writing program this way programmers can solve problems by breaking them up in to smaller more manageable parts. Modularity furthermore allows parts of the program to be reasoned about independently. Python has several mechanisms to achieve modularity, it has functions, classes, modules and packages. This chapter will explain how functions and modules can be used to create a modular program. Reusability is also enabled by modularity because the independent parts of a program can be included in other programs.

While programs are executed by computers they are written, adapted, and maintained by humans which implies that source code must be written to be read *and* understood by humans. Readability of code is affected by many factors; clear structure, descriptive naming schemes, adhering to language idioms, useful comments, clear documentation, and so on. This chapter will explain how Python code can be documented and give some recommendations for naming schemes.

8.1 Functions

Functions are a way to name a block of executable code, much like a variables are a way of naming values, and like variables allow you to reuse a value, functions allow you to reuse a blocks of code. They take a number of inputs, perform some operations and return some outputs. The input variables that appear in the function definition are called the *parameters* of that function, while the actual values that are sent to the function are called *arguments*. A function definition starts with the def keyword, followed by the name of that function, a list of parameters enclosed by parentheses, a colon and finally an indented block of code called the *function body*. In Python you do not need to declare the data type of the parameters to a function. The first example shows two simple function definitions:

```python
def say_hello():
    '''Print 'Hello!'.'''
    print 'Hello!'


def is_even(n):
    '''Return True if n is a even.'''
    return n % 2 == 0
```

The say_hello function in this example takes no arguments (since there is no list of arguments between parentheses) and it has no return values[1] (since there is no return statement). When you use the first function it will print Hello! to the standard output. The second function is_even takes only one argument, a number n, and returns True or

---

[1]Beware that Python considers the output of a function without a return statement to be None.

False depending on whether the number is even or not. Both functions contain a triple quoted string with an explanation of what the function does. These types of strings, so-called docstrings, will be explained in Section 8.3. The next example shows how you can filter a list of numbers so that the resulting list only contains even numbers:

```python
numbers = [1, 2, 3, 4, 5, 6, 7, 8]


def is_even(n):
    '''Return True if n is a even.'''
    return n % 2 == 0


def filter_even_numbers(l):
    '''Return a list of even numbers in l.'''
    out = []

    for number in l:
        if is_even(number):
            out.append(number)
    return out


print 'The even numbers are:', filter_even_numbers(numbers)


print out  # This will fail with an NameError because out is unknown here.
```

The final line of this little program demonstrates that functions can define variables that are invisible to the rest of the program. This allows you to reuse obvious variable names in different contexts without having to worry that you already used them somewhere else. The following section will explain the rules around visibility and how you can use these rules to create a nice modular program.

8.2   Scope

The NameError in the last example of the previous section showed that variables can be defined inside of a function that remain invisible to the rest of a program. The extent of program where the variable is visible is called its *scope*. In Python functions, classes and modules all create their own scopes. These scopes can be nested, and the outermost scope is called the global scope. Names defined in the global scope are visible throughout the whole program — even inside of functions that have their own scope! When a name is accessed it is first searched in the local scope, then in one (or more) enclosing scopes until the global scope is reached. When the name cannot be found in the global scope either, the Python built-ins will be searched. A failure to find the name in the built-ins will finally lead to a NameError (as you saw in the last example of the previous section).

Variables defined in the global scope are accessible everywhere. While it may seem appealing to have all your variables in the global scope so that you can access them everywhere it is a bad idea, really! In a large program that uses many global variables every part of the program has the potential to affect every other part of the program, and therefore the program will be hard to understand, debug or adapt.

When a name that exists in the global scope is assigned to from a enclosed scope , that object in the global scope is not changed. In stead, the name is defined in the local scope making the global object inaccessible from the local scope. To circumvent this problem Python has the global keyword as was demonstrated in the example above. If you are only interested in the value of a global without needing to change it you can leave the global keyword out. Trying to change the value of a global variable without using the global keyword will result in a UnboundLocalError. The **global** keyword is rarely needed, and should be avoided. The following example shows the behavior of global variables.

```python
a = 10  # in global scope


def add(a, b):
    # Note: a and b are now local variables of this function
```

```python
    # Note: the local a variable is something else than the global one
    a = a + 10  # all working on local variables!
    return a + b

def change_a():
    # To change the value of a global variable from some function's scope
    # use the global keyword of Python.
    global a  # Is needed!
    a = a + 10

def two_a():
    # Note: you can access a global variable to read its value without the
    # global keyword.
    return 2 * a

print a  # prints 10
print add(2, 2)  # prints 14
change_a()
print a  # prints 20
print two_a()  # prints 40
```

8.2.1   Modules

A Python file is called a module. Modules can be run as programs (as was demonstrated in Section 3.2 or imported. To import a module means to load its definitions so that they may be used in the program you are working on (or in a Python shell). The load for instance the math module that contains more advanced mathematical functions than what is normally available use the import math statement. This will create a name math in your program (or Python shell) that points to the math module. It is also possible to import a module under a different name, or just some functions from a module:

```python
import math
import cProfile as profile  # imports the cProfile module and names it profile
from os.path import join  # imports only the join function from os.path

# and now a really bad and lazy idea !!!
from pylab import *  # import all definitions from pylab into the global scope
```

The last import statement of the above example is a really bad idea because it indiscriminately pulls all definitions in the pylab module into your program (which may overwrite some of the other definitions). You should only ever import those things you need, or a whole module in its own name space (that is give it a name and access all other functions via that name using the . notation). To access the cosine function from the math module in this example using the unambiguous math.cos notation.

8.3   Documentation

Python code can be documented in several ways. Code with a logical flow, clear modularity, and well chose, descriptive, names is referred to as self documenting (the holy grail of well written code). Comments, i.e. lines starting with a #, are used to explain the purpose of sections of code, the reasons it was written a certain way, or to give some extra context. Python also allows so-called docstrings which are triple quoted strings at the top of a module, function, class, or class method that explain what their purpose is. Docstrings can easily be turned into nice documentation using one of Python's documentation tools[2]. Docstrings should explain how to use the module/class/function/method, without a need to read the underlying code.

---

[2]See, for instance, http://sphinx-doc.org/.

The following example shows a simple module with a module docstring and two function definitions with their own docstrings (for the purposes of this text assume that this module is named `simple.y`).

```python
#!/usr/bin/env python
'''
Some trivial string manipulations.
'''


def is_palindrome(word):
    '''
    Check whether a word is a palindrome.
    '''
    # the [::-1] slice goes through a string in reverse
    return word == word[::-1]


def reverse_case(word):
    '''
    Reverse the case of individual letters in a word.
    '''
    # create a list with the case reversed letters:
    case_reversed = []
    for c in word:
        if c.isupper():
            case_reversed.append(c.lower())
        elif c.islower():
            case_reversed.append(c.upper())
        else:
            case_reversed.append(c)

    # turn the case reversed list in string and return it:
    return ''.join(case_reversed)
```

You can import a module in an interactive Python session, and use the built-in `help` function to view that module's documentation. The documentation will appear in a viewer as shown in Figure 8.1.

```python
>>> import simple
>>> help(simple)
```

8.4 The `if __name__ == '__main__'`: trick

Python modules can be used in two contexts. They can be executed directly or they can be imported by other Python modules. The only way of differentiating these two cases is to check a special variable __name__ whose value will differ based on the context. If a module is executed directly the value of that variable will be the string '__main__'. If the module is imported its value will be the name of the module. This variable is automatically set by the Python interpreter.
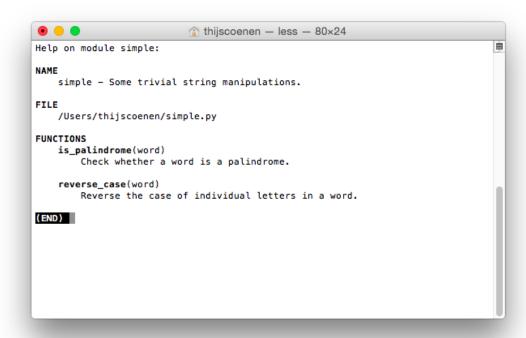
Figure 8.1: The Python built-in `help` function in action viewing the module documentation of the `simple.py` module.

Chapter 9

Matplotlib

9.1 Introduction

Matplotlib is the standard when it comes to plotting Python. There are other packages available, but none is as full featured and widely used as Matplotlib. Matplotlib was inspired by Matlab's plotting facilities, which at times makes its API (Application Programming Interface) feel a bit foreign to an experienced Python programmer. What follows are a few examples of standard plotting tasks in Matplotlib to get you started. Extensive documentation is available online. One note about older Matplotlib documentation, some of it uses `pylab`, which you should avoid as it makes scripts hard to maintain (look for examples using `pyplot` in stead).

9.2 Simple use of `pyplot`

You can plot in several ways using Matplotlib. The simplest way is to use the `pyplot` module and allow it to keep track of which figure and plot you are working on. The first two code examples demonstrate this way of using `pyplot` (screenshots are shown in Figure 9.1). Matplotlib includes a built-in viewer, that can be invoked using the `pyplot.show()` function. This viewer can be used to zoom in on a plot and to save it. Below the code that creates the scatter plot using the `pyplot.scatter` function. It expects one or two lists (or arrays) of coordinates, one for the x-coordinates and one for the y-coordinates, these lists (or arrays) must be of the same length.
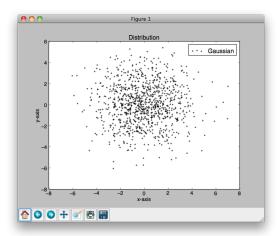
```python
from matplotlib import pyplot
from random import normalvariate


# generate data:
X1 = [normalvariate(0, 2) for i in range(1000)]
Y1 = [normalvariate(0, 2) for i in range(1000)]


# create plot:
pyplot.scatter(X1, Y1, color='k', marker='+', label='Gaussian')
pyplot.legend()
pyplot.title('Distribution')
pyplot.xlabel('x-axis')
pyplot.ylabel('y-axis')
pyplot.show()
```
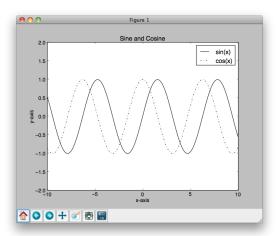
The `marker` keyword argument of `pyplot.scatter` sets the marker to be used for each point. Possible markers are o for circles, . for dots, + for plusses, x for crosses, * for stars and several other possibilities that can be found in the Matplotlib documentation. The color of the points is set using the c keyword argument, you can pass it any color that Matplotlib understands[1]. The `pyplot.title` function sets the title for the entire image and `pyplot.xlabel`

---

[1]`http://matplotlib.org/api/colors_api.html`

(a) Simple scatter plot using `pyplot.scatter`.   (b) Simple line plot using `pyplot.plot`.

Figure 9.1: These two screenshots show the interactive viewer included with Matplotlib. The interactive viewer allows you to zoom in on the data and save the plots as well.

and `pyplot.ylabel` set respectively the label for the horizontal axis and for the vertical axis. The legend was created by labeling the data that was plot (the `label` keyword argument of `pyplot.scatter` and subsequently using the `pyplot.legend` function.

The second example (below) shows how line plots can be constructed using the `pyplot.plot` function. Just like `pyplot.scatter` this function expects one or two lists (or arrays) of numbers to be used as coordinates. The color and style of each line can be set in several ways but here we have used the shorthand notation for a black line and a black dot-dashed line, respectively `'k-'` and `k-`. Note that this example uses Numpy arrays, and functions to create the data.

```python
import numpy
from matplotlib import pyplot

x = numpy.linspace(-10, 10, 201)

pyplot.plot(x, numpy.sin(x), 'k-', label='sin(x)')
pyplot.plot(x, numpy.cos(x), 'k-.', label='cos(x)')
pyplot.ylim([-2, 2])
pyplot.xlabel('x-axis')
pyplot.ylabel('y-axis')
pyplot.legend()
pyplot.title('Sine and Cosine')
pyplot.show()
```

Axes can be limited to a certain range using `pyplot.xlim` and `pyplot.ylim` functions to respectively constrain the horizontal and vertical axes.

The third example shows a simple historgram created using the `pyplot.hist` function. Unlike the plots in the earlier examples, which were shown using the interactive plotting component of Matplotlib, the third example saves the figure directly to a PDF (Portable Document Format) file. Using PDF is a good idea because Matplotlib saves your plot as a vector image which will look sharp and clean even when enlarged.

```python
import numpy
x = numpy.random.normal(5, 1, 10000)
```

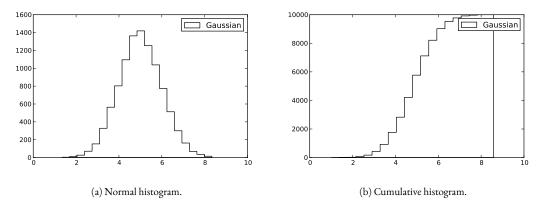(a) Normal histogram.   (b) Cumulative histogram.

Figure 9.2: Simple histograms created with `pyplot.hist` and saved as a PDF file using `pyplot.savefig`. The cumulative plot on the right was created using the <u>`cumulative=True`</u> keyword argument to the `pyplot.hist` function.

```
pyplot.hist(x, 20, histtype='step', color='k', label='Gaussian')
pyplot.xlim(0, 10)
pyplot.legend()
pyplot.savefig('pyplot-histogram.pdf')
```

The first argument to `pyplot.hist` is a list or array of number for which a histogram is constructed. The second, optional, argument is the number of bins that should be created (20 in this example). The `histtype` keyword argument controls how the historgram looks and can be set to `bar`, `barstacked`, `stepfilled` or `step`. The color of the histogram is further controlled using the `color` keyword argument. The histogram can be turned into a cumulative histogram setting the `cumulative` keyword argument of `pyplot.hist` to `True`. Similarly the `normed` keyword argument can be set to `True` to normalize the histogram such that the histogram will sum to 1. The `orientation` keyword argument that defaults to `'vertical'` argument controls the direction of the bars in the histogram.

As is visible in Figure 9.2 the default position of the legend can be inappropriate. The `loc` keyword argument to the `pyplot.legend` function controls the position of the legend in the figure and can take the following values: `'best'` (which is default), `'upper right'`, `'upper left'`, `'lower left'`, `'lower right'`, `'right'`, `'center left'`, `'center right'`, `'lower center'`, `'upper center'`, and `'center'`.

9.3 Multiple plots in one figure

To add several plots to one figure recent versions of Matplotlib have the `pyplot.subplots` function to create an array of subplots while older versions have the less convenient `pyplot.subplot`.

9.3.1 Using **pyplot.subplots**

The following example shows how to create a figure with two subplots in it.

```
import numpy
from matplotlib import pyplot

fig, axes = pyplot.subplots(nrows=1, ncols=2)
X = numpy.linspace(-10, 10, 201)

axes[0].plot(X, numpy.sin(X), 'k-', label='Sin(x)')
axes[0].legend()
axes[0].set_title('Sine')
```
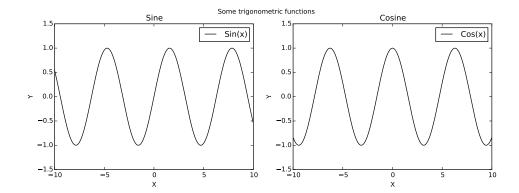
Figure 9.3: Two subplots created using `pyplot.subplots`.

```
axes[0].set_ylabel('Y')
axes[0].set_xlabel('X')
axes[0].set_ylim([-1.5, 1.5])

axes[1].plot(X, numpy.cos(X), 'k-', label='Cos(x)')
axes[1].legend()
axes[1].set_title('Cosine')
axes[1].set_ylabel('Y ')
axes[1].set_xlabel('X')
axes[1].set_ylim([-1.5, 1.5])

w, h = fig.get_size_inches()
fig.set_size_inches(2 * w, h)
fig.suptitle('Some trigonometric functions')
fig.savefig('subplots-example-1.pdf')
```

The `fig` object a reference to the figure as a whole (an instance of `matplotlib.figure.Figure`) and the `axes` object is an array of `matplotlib.axes._subplots.AxesSubplot` instances. To set the limits and labels on the individual plots you can no longer use the `xlabel`, `ylabel`, `xlim` and `ylim` functions of `pyplot`, instead use the `set_xlabel`, `set_ylabel`, `set_xlim` and `set_ylim` methods of the individual plots. The example furthermore demonstrates how the size of the figure can be found and changed using the `get_size_inches` and `set_size_inches` methods of the figure. If you want to set an overall title for the figure use the `suptitle` method of the figure object. It should be noted that `pyplot.subplots` will return a two-dimensional array if both `nrows` and `ncols` are larger than one, a one-dimensional array if either `nrows` or `ncols` is one or just a reference to a subplot if both `nrows` and `ncols` are one.

9.3.2   Using `pyplot.subplot`

While `pyplot.subplots` returns subplot and figure references, `pyplot.subplot` does not, in stead it activates one of the subplots and all subsequent drawing commands go to that subplot. Below a simple example showing a scatter plot and two one-dimensional histogram (see also Figure 9.4).

```
from matplotlib import pyplot

X = numpy.random.normal(0, 1, 1000)
Y = numpy.random.normal(0, 1, 1000)

pyplot.suptitle('Some two-dimensional distribution')
pyplot.subplot(2, 2, 1)
pyplot.hist(X, 10, orientation='vertical', histtype='step', color='k')
```
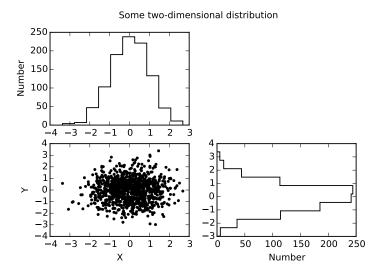
Figure 9.4: Three subplots created using `pyplot.subplot`.

```
pyplot.ylabel('Number')
pyplot.subplot(2, 2, 3)
pyplot.scatter(X, Y, marker='.', color='k')
pyplot.xlabel('X')
pyplot.ylabel('Y')
pyplot.subplot(2, 2, 4)
pyplot.hist(Y, 10, orientation='horizontal', histtype='step', color='k')
pyplot.xlabel('Number')

pyplot.savefig('subplot-example-1.pdf')
```

The arguments to `pyplot.subplot` are: the number of rows, the number of columns and *plotnumber* of the subplot to be activated. The *plot number* numbering scheme starts with the top left plot and moves through the other plots in that row before moving to the next row, etc. These arguments can be supplied separately as is done in the example above, or a shorthand can be used in special circumstances. Provided that each of the parameters to `pyplot.subplot` is smaller than 10, the numbers can be combined into one parameter. To activate the top left plot in an table of plots with two rows and two columns you can use the shorthand 221, where each position in the number will be treated as one of the arguments — the call then becomes `pyplot.subplot(211)`. This shorthand is used quite often in examples in the Matplotlib documentation (but is not very *Pythonic*).

9.4 Customizing tickmarks

While Matplotlib conveniently creates tickmarks in appropriate locations for your figures, you can also control them yourself using `pyplot.xticks` and `pyplot.yticks`. The following example shows how the position, text, and rotation of tickmarks and their labels can be controlled (see Figure 9.5). For more precise control see `matplotlib.ticker` in the Matplotlib documentation.

```
import calendar
from matplotlib import pyplot

months = list(calendar.month_name)[1:]
temps = [3.4, 3.5, 6.1, 9.1, 12.9, 15.4, 17.6, 17.5, 14.7, 11.0, 7.1, 4.0]

pyplot.scatter([i for i in range(12)], temps, marker='x', color='k')
```
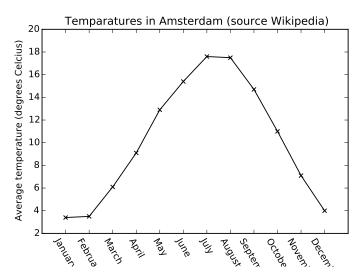
Figure 9.5: Controlling tickmark label position, text and rotation.

```python
pyplot.plot([i for i in range(12)], temps, 'k-')
pyplot.xticks([i for i in range(12)], months, rotation=-60)
pyplot.xlim([-1, 12])
pyplot.title('Temparatures in Amsterdam (source Wikipedia)')
pyplot.ylabel('Average temperature (degrees Celcius)')
pyplot.savefig('xticks-1.pdf')
```

Chapter 10

Further reading

10.1 Course goals

The goals for the Basic Linux and Coding for Astronomy & Astrophysics course have been described in the Introduction to this reader. Because all of Python cannot be learned from scratch in the amount of time allowed for the course, a practical subset of Python was defined that you must understand to pass. The followinf subsections list the parts of Python you are expected to understand. In this context "understand" means that you can explain them in your own words and that you can apply that knowledge when designing and implementing programs.

10.2 A practical subset of Python

You are expected to know the basic Python data types, what they are used for, how to convert into each other, how to iterate over the various sequence and mapping types (both uing indices and by direct iteration) and how to perform simple operations on them (including arithmetic). Furthermore you should know how you find the type of some value or variable.

- the assignment operator

- number types: floating point numbers and integers

- sequence types: strings, tuples, and lists

- mapping types: dictionaries

- Booleans and `None` type

- the literal notations for these data types (!)

- indexing and slicing of strings, tuples, and lists

- immutable versus mutable data types

- the `type` function

- the `int`, `float`, `str`, `list`, `dict`, and `tuple` functions

- `list.sort`, `list.append`, and `list.pop` methods

- `min`, `max` and `len` functions

You are expected to understand Boolean expressions, Boolean operators and comparison operators.

- `and`, `or` and `not`

- Truthy and Falsy as applied to strings, tuples, lists and dictionaries

- the comparison operators (like == etc.)

- the `bool` function

- the `in` operator as applied to strings, tuples, lists and dictionaries

You should know how to conditionally or repeatedly execute a blocks of code. You should also know the different ways you can iterate over strings, tuples, lists, and dictionaries.

- `if/elif/else`

- `for/else` and `while/else`

- the `break` statement

- the `continue` statement

- the `range` function

- the `enumerate` function

You should be able to write a properly structured Python program. Meaning that the program is modular — for the purposes of this course modular at the level of functions and modules is enough (you are not expected to understand classes and packages are not necessary). While designing a program you should be able to break up a problem into smaller parts that are more easily solved. The functions you implement should have clear purposes and be as indepentent of the rest of the program as possible — meaning do not use the global scope to communicate state between functions, make sure functions only solve one part of the problem. You should understand what scope is, how scopes nest, which Python constructs create scopes of their own and which do not. Your modules should allow reuse of their contents so you should understand how you can prevent code from running at import of a module. Furthermore you should be able to document your code and know the difference between documentation and comments in your code.

- the `def` keyword and function definition

- what scope is and how it is used

- the `if __name__ == '__main__':` trick

- docstrings (for modules and for functions) and how they differ from comments

Last but not least you should be able to write programs in proper Python style, which includes proper use of white space, descriptive naming schemes, comments, and indentation (4 spaces). In general you should follow the Python style rules as laid out in PEP 8.

10.2.1   Numpy and Matplotlib

You are expected to be able to create plots and use Numpy arrays during the homework exercises. You will not be expected to know them by heart during the exam. You are, however, expected to be able to explain at a high level the use of arrays and what makes them different from lists.

10.3   More Python resourses

Python has a large community that produces a large number of libraries and a large amount of documentation. Various learning resources are available for free on-line.

- For the absolute beginner there is the *Learn Python the hard way course* by Zed Shaw. It assumes very little computer knowledge and takes small steps in teaching.

- As an alternative to the Python parts of this reader, you can also read the relevant chapters of *Think Python* by Allen B. Downey. This book is an appropriate resource for learning Python 2. It is a more complete introduction to Python than can be given in this course.

- If you already know quite a bit of Python the *Intermediate Python* website is a nice resource to dive deeper in more advanced Python constructs.

Part II

Assignments

Chapter 11

Finding Pulsars from the command line

11.1   Introduction

Pulsar surveys in radio produce large numbers of potential pulsars, these then have to be inspected. For this exercise you are given a large number of potential pulsars. For each potential pulsar you will find a `.bestprof`-file and a `.png`-file. The former is a text file with meta data while the latter is a diagnostic plot. Below I will use detection to describe one of these potential pulsar signals (whether a real pulsar or not). The data set should contain a pulsar that you have to find using command line tools.
The data you need is available from the `http://astro.mrprog.nl` website.

11.1.1   Grading

This exercise is required but will be graded PASS or FAIL. Hand in your work via the `http://astro.mprog.nl` website.

11.1.2   The assignments

- Start a report to be handed in as a A4-sized PDF file. Make sure you include your name and student number on the first page of your report.

- Download the file `LOTAAS.tar.gz` (containing a number of pulsar non-detections and pulsar detections) and copy it to a directory where you can work.

- Unpack the tarball, you will now be presented with a number of sub directories containing the images and metadata files.

- Open one of these `.png` images from the command line. If you don't know which command to use to open these images, perform a web search to find out which command will do that. (Add the command you used to the report.)

- Now that you have seen the plot, open the metadata (hint: the bestprof files are text files). Looking through this file you will see that the reduced chi-squared is an entry in the header, we will need these later. How many entries are there in the header? (Add the number to your report.)

- How many possible detections are there in the complete data set? (Add the answer to your report.)

- Add the picture of the brightest profile to your report (the brightest detection will have the highest reduced chi squared). You should use the meta data to find the brightest non-detection. Also add to the report how you found the brightest one.

- Now add the 100th brightest profile you report and explain how you found it. (The reader contains a hint as to how you could do this.)

Chapter 12

Berkeley Earth and data in text files

Berkeley Earth[1] is a project that provides an independent reconstruction of Earth's recent temperature history. It took several sets of historical weather station observations and used them to perform the reconstruction. All data that Berkeley Earth used, both raw and transformed (filtered for bad data, adjusted for systematics etc.) are freely available from the project's website. We will be using Berkeley Earth's data for this exercise. We will not perform careful analysis, but use the data to gain experience reading and writing textual data sets.
In this exercise we will:

- read and write files,

- create loops, and

- learn to use the basic Python data types.

12.1   Assignment

For this exercise you cannot use any modules not in the Python standard library, and furthermore you are also not allowed to use the `csv` module that is available from the Python standard library. Start a report (A4, PDF) that you use to explain the steps you took. This report is to be handed in along with the Python scripts you will create, use it to answer questions below and if needed to explain your methods.

12.1.1   Reading data

The first step is to write some Python code to process a single file of Berkeley Earth data.

- Go to the data section of Berkeley Earth's website and read the data overview before doing anything else.

- Download the monthly average temperatures for land only. You will write a simple Python program to read these data. Open the data file in a text editor and have a look at its content. You will see that there is a header containing some meta data and after that columns of data. Note that the data file contains the temperature anomaly, not the actual temperature.

- Create a new Python script called `datareader.py`. Be sure to add a comment, i.e. a line starting with a `#`, with your name and student number to the file — Python will ignore this line. Make sure that the data file is in the same directory as you Python program (for convenience). To confirm you can execute the program make it print hello world to the terminal.

- Open the data file using the `open` function as described in Chapter 7.

- Now open the data file and iterate over its contents while printing each line to the terminal.

---

[1]See `http://berkeleyearth.org/`.

- You are now ready to write a loop that extracts the actual data from the data file. The format for the data that you should make is a list of tuples, where each tuple contains the year, the month, the temperature anomaly and its uncertainty.

- For your report: which data types in Python are appropriate to represent the aforementioned values?

- Write a loop that extracts the aforementioned values and creates the list of tuples containing the values. Be sure to choose an appropriate value to represent missing data (the NaN values in the data file).

- For extra credit: extract the January 1951 - December 1980 temperature values from the file and add them to the temperature anomalies you extracted before to arrive at real temperatures (all using Python and an unmodified data file).

- For your report: which month and year has the largest temperature anomaly and which the smallest (write some code to do this)? Report also the values of either the temperature anomaly or the temperature.

- To be handed in: Your Python code saved as a file `datareader.py`.

12.1.2   Refactoring

Choose any 4 weather stations with at least 250 measurements each available from the Berkeley Earth data set. You will now write a small program to process these files in such a way that one program can process all 4 files.

- Create a new Python script and as before create a loop to extract the year, month, temperature anomaly (or temperature anomaly + average) and uncertainty. Do this for one data file.

- We are now going to use the template provided in Section 12.1.5 to write a simple program that can process each of the data files you choose. You should copy the template to a file called `stationdata.py`. The template contains a function definition (the line starting with `def`) that defines a function `read_file` that accepts one parameter, the name of the file it operates on) that you can use as is. Before changing program confirm that it runs without crashing.

- Adapt your code so that it works inside of the function `read_file`. If you did this correctly the program should print out how many data entries each file has.

The process of re-organizing code to be better organized or a better fit to the problem at hand is called *refactoring*. When you write larger programs you will often have to refactor parts of them. In this case you made a piece of code more general, going from being able to process only one file, to being able to process an arbitary number of files — assuming their contents are laid out in the same way. Next week we will treat function definitions and program structure in more detail.

12.1.3   Writing data

You will now write a small program to write the anual average temperature anomalies to a text file. You can re-use parts of the programs you wrote earlier.

- Create a new Python file called `anualaverages.py`.

- Note that for this exercise you do not need to define functions yourself yet — a "structureless" program is ok for now.

- Copy the file reading code from either of the previous exercises and extend it with a way of calculating anual averages. Make sure that you take into account months that have missing data properly in your averages.

- Make sure that the format of the data you produce is a list of tuples where each tuple contains the year, temperature (anomaly) and uncertainty on that temperature (anomaly) represented by appropriate data values.

- Now add code to open a file for writing, and loop over your anual averages. For each entry write a line to the file with a column for the year, a column for the temperature (anomaly) and a column for the uncertainty on that temperature (anomaly). The columns should be separated by white space (a tab or a space).

- You are now almost ready, but the data file you wrote has no header describing what it contains. Add a code to write a line explaining what each column contains. These also should be white space separated and easy to read with by a program.

12.1.4   Hints

- To represent a newline Python uses the \n escape sequence, lines accessed in a file will end with this character on UNIX. The '\n' character is also called *line feed*. On Windows the end of lines need '\r\n', a combination of *carriage return* and line feed. You will need to strip these characters when processing the files (look at the strip() method of strings).

- Look up the split, join, startswith and endswith methods of strings on the Python website — you will need them.

- Look up the min and max functions.

12.1.5   Program template

```python
#!/usr/bin/env python
'''
Small program to read single weather station data from Berkeley Earth's data set.
'''


def read_file(filename):
    '''
    Read a single data file given by filename.
    '''
    data = []  # prepare an empty list that you will fill
    # --------------------------------
    # Add your file handling code here:

    # --------------------------------
    return data  # return the data to

if __name__ == '__main__':
    filenames = [
        'YOUR FIRST FILENAME HERE',
        'YOUR SECOND FILENAME HERE',
        'YOUR THIRD FILENAME HERE',
        'YOUR FOURTH FILENAME HERE',
    ]

    for filename in filenames:
        data = read_file(filename)
        # You need not understand the following line for now:
        print 'File %s contains %d data entries' % (filename, len(data))
```

12.2   Submitting

The submission for this week should contain:

- A report (A4, PDF) called `week3.pdf`

- A Python program called `datareader.py`

- A Python program called `stationdata.py`

- A Python program called `anualaverages.py`

- A `gzip`-ed tar ball `stationdata.tar.gz` with the data file you used for the second program.

Chapter 13

Plotting Berkeley Earth data (Group A, week 4)

In this week's assignment you will create plots of more Berkeley Earth data. The goals for this week are to:

- write a small program split into several modules,

- create properly structured modules,

- document code properly, and

- create the basic plots (scatter plot, line plot, histogram).

Note: properly label your plots also have a look at the hints section below.

13.1 Gathering data

Like last week we will be using Berkeley Earth data because it is nice and homogeneous. The course website has a tarball with all the necessary files.

- Download the weather data provided in a tarball on the course website `astro.mprog.nl` use one of the files in that tar ball to check that the parser you write works.

- Create a Python module called `weather_parser.py`. In this module create a function called `parse_location_data` that accepts a file name (as string) and returns a list of tuples with each tuple containing a year (as integer), a month (as integer), a temperature anomaly (as a floating point number), and the uncertainty on the temperature anomaly (again a floating point number. It is a good idea to refer to the code you wrote last week.

- Add a `if __name__ == '__main__':` line to the module and move any code that calls the `parse_location_data` function to the block behind it. This will ensure that at import your module will not start parsing a file — in other words, that your module is safe for re-use.

13.2 Creating a line plot

You will now start creating plots using Matplotlib.

- Create a Python module called `plotter.py`, you will use this module for all the plotting tasks below. Import the `pyplot` module of Matplotlib using: `from matplotlib import pyplot`

- Import your `weather_parser` module.

- Create a `main` function that you will use to call the various plotting functions.

- Create a function called `plot_timeseries` that uses that `pyplot.plot` function to create a line plot of the temperature anomaly for one of the data files. The plot should be saved as `time-series.pdf`. As a X coordinate you are allowed to use years as floating point numbers (with the fractional part set to the appropriate value for each month) if you think using Python dates is too hard.

- Call the `plot_timeseries` function from `main` function.

- Extra credit (do this after you finish the rest of the assignment): Plot a gray background around the time-series showing the uncertainty in the data (like the plots on the Berkeley Earth website have).

13.3 Creating a scatter plot

- In your `plotter.py` module create a function `plot_scatter` that creates a scatter plot of uncertainty versus year that the data was taken. Is there any pattern you see? Call this function from `main` and save the resulting plot as `scatter.pdf`

13.4 Creating a histogram

For this part of the exercise you should download the data for all the locations in the data file.

- Create a function called `plot_histogram` in `weather_parser.py` that accepts a list of file names (for the location data files) and produces a histogram of the temperature anomalies of all the cities at January 2000.

- You will need to loop over all the file names and call the parser function for each location. Keep a list of all the temperature anomalies, use that to make the histogram.

- Save the histogram as: `histogram.pdf`

13.5 Hints

- For a plot of two variables `pyplot.plot` expects two lists of equal lengths, one containing all X values and one containing all Y values. In our case that means that you have to take the list of tuples you created and turn it into separate lists (for time values, for temperature anomalies etc.).

- If you are having problems importing your own module make sure that both modules (the one doing the importing and the one being imported) are both in the same directory.

13.6 Handing in

You should make a gzipped tarball (called `week4.tar.gz`) with the following files:

- `weather_parser.py`

- `plotter.py`

- `time-series.pdf`

- `scatter.pdf`

- `histogram.pdf`

- `report-week-4.pdf` (Only if you want to explain your methods.)

Chapter 14

Plotting LOTAAS survey results (Group B, week 4)

In this exercise you will take the LOTAAS data set used in week 3 to create some plots. Unlike some other data sets that you may encounter the data are spread over a few hundred small files spread over a directory tree. The goal for this exercise is to write a program that gathers the necessary data and creates plots for these data. Goals for this week are to:

- write a small program split into several modules,

- create properly structured modules,

- document code properly, and

- create the basic plots (scatter plot, line plot, histogram).

Note: properly label your plots also have a look at the hints section below.

14.1 Parsing **.bestprof** files

As you have seen before the .bestprof files contain easily extracted data about each pulsar candidate. Before we start creating plots we first have to write a parser for these data files. As last week, use should write the parser yourself. Do not use Pandas, Numpy or Scipy data reading functions or `csv` module.

- Create a Python module called `bestprof.py`, this is where the .bestprof file handling will go.

- Have a look through a .bestprof file with a text editor. Note that the file consists of two parts, a header with key-value pairs and a simple two-column list of values. The latter is the pulse profile split into a number of bins (how many is a parameter chosen by those who run the LOTAAS survey).

- Create a function `parse_bestprof` that accepts a filename (as a string) and that returns a dictionary for the key-value part of the header and a list of values for the list part of the .bestprof file (i.e. for the pulse profile). *Note: Python allows multiple returns from functions.*

- Note: turn the interesting parts of the header into the appropriate data types. At least do this for the Dispersion Measure DM, the pulse period P, pulse period derivative $\dot{P}$ and reduced chi-squared χ_r. Do the same thing for all values in the pulse profile.

- Extra credit: split the file parsing into two functions, `parse_header` and `parse_profile`, that get called by `parse_bestprof`.

14.2 Walking a directory tree

Now that you are able to parse individual `.bestprof files` you must must write some code to be able to find all relevant files.

- Create a Python module called `finder.py` that will contain the code that searches for `.bestprof` files.

- Look up the function `os.wallk`, as it will allow you to search for file much like the BASH `find` command does.

- Write a function called `find_bestprof` that takes a directory (as a string) and returns a list of `.bestprof` files names. *This involves writing a loop over all subdirectories using the `os.walk` function, checking the filenames and accumulating the relevant filenames in a list that you return.* Make sure that you return not just sub directories but the full paths to the files.

- *You can check your code by comparing the output of the BASH `find` command with the output of your code. Do both find the same number of files?*

14.3 Plotting survey results

You will now tie together all the code you wrote before and create plots for the whole set of survey data.

- Create a Python module `lotaas_plotter.py` and make sure that it is saved in the same directory as the `bestprof.py` and `finder.py` modules.

- Create a `main` function function that accepts a directory name (as string) and that will contain all functionality of this program.

- Confirm that you can `import` the functions defined in the modules you created earlier on. Fix it if this does not work.

- Call the `find_bestprof` function to gather a list of `.bestprof` filenames. Loop over this list and call the `parse_bestprof` for each of the file names. You should create a list of tuples for the output of all these functions.

- Below are three plotting functions that should all be called from the `main` function.

- Create a function `plot_p_pdot` that uses Matplotlib to create a scatter plot of logarithm of P versus the logarithm of $\dot{P}$. You are allowed to throw away values for which you cannot take the logarithm. If you are unsure about what these plots look like use the internet to find some examples of P-Pdot diagrams. Save the resulting plot in a file called `ppdot.pdf`.

- Create a function function `plot_profile` that creates a line plot of the pulse profile value versus bin number. Make sure to include the period and reduced chi-squared values in the title of the plot. Save the profile of the brightest signal (or of the pulsar signal if you can find it) to a file called `pulseprofile.pdf`. Extra credit: create a plot of value versus time (you will need to grab a number from the header to make this work).

- Create a function called `plot_histogram_p` that creates a histogram of all the pulse periods in the suvey data. Save the resulting plot to a file called `p_histogram.pdf`.

- If you did not do so already: Add a `if __name__ == '__main__':` line to your program and call the main function from the block of code behind it.

- Using the `sys.argv` variable it is possible to create a script that accepts command line arguments — but doing so, is the quick and dirty way. If you want you can change your `lotaas_plotter.py` to accept the directory where it needs to search as a command line argument. *The "proper" way of creating command line arguments is to use a module like `argparse`.*

14.4   Hints

- If you are working on university computers, copy your data to the scratch disk in a directory with your name. Work in that directory and only save your scripts to your home directory (because the network is quite slow).

- To do path manipulations in Python (like creating or concatenating full paths) look at the `os.path` module of Python. Interesting functions in that module are `os.path.join`, `os.path.split`, `os.path.abspath`, and `os.path.realpath`.

- If you are seeing old plots when you are creating new ones with Matplotlib you should probably use the `matplotlib.pyplot.clf` function to clear the old figures.

- A good rule of thumb for the number of bins in a histogram is the square root of the number of items in the data.

14.5   Handing in

You should make a gzipped tarball (called `week4.tar.gz`) with the following files:

- `bestprof.py`

- `finder.py`

- `lotaas_plotter.py`

- `ppdot.pdf`

- `pulseprofile.pdf`

- `p_histogram.pdf`

- `report-week-4.pdf` (Only if you want to explain your methods.)

Chapter 15

Numpy arrays and image manipulation

For this week's assignment you will do some simple image manipulations and plots of two-dimensional arrays. The goals for this week are to:

- create and resize Numpy arrays,

- introspect size and data types,

- use indexing and slicing for arrays,

- use `imshow` to show two-dimensional arrays,

- do some simple array manipulations.

Be sure to stick to the PEP 8 style rules, and do not forget to properly label any plots you create.

15.1   Creating a simple one-dimensional plot

- For this part of the assignment create a Python script called `planck.py`.

- Make a plot of the $B_\lambda(\lambda, T)$ Planck curve [1]

- Use `numpy.linspace` to create an array of wavelength values.

- Implement Planck equation in terms of `numpy` math functions that operate on arrays.

- Add the plot to your report.

15.2   Image manipulation

15.2.1   Loading the image

- Download the image provided on the course website, you will use this image throughout the assignment. Create a new Python script and call it `imagemanipulations.py`, in that script create a function `transpose_image` where you add this section's code (to test call that function). This function takes one argument, the name of the file that should be transposed.

- As a first step you are going to load the image and turn it into a Numpy array. There are several functions you could use to perform this step, but we will use `scipy.ndimage.imread`. Look up the documentation for this function and load the image into an array.

- For your report: what are the dimensions of this image and what are the data types in the array?

---

[1]See for instance `https://en.wikipedia.org/wiki/Planck's_law`.

- Use `pyplot.imshow` to show the image using Python. Transpose the X and Y axes of the image using the `numpy.swapaxes` function. Include the transposed image in your report.

15.2.2   Making a color separation

Images as displayed on your screen can be separated in a red, a green, and a blue channel.

- Start by creating function `color_separation` for this section of the assignment. This function should take one argument, the filename of the image.

- Load the image into an array.

- Create a grid of four sub plots using the `pyplot.subplots` function. In the top left plot put the original image, use the other sub plots to show the individual color channels.

- The image consists of a red, a green and a blue color channel. Use array indexing / slicing to create arrays for the individual color channels. Add the individual channels to the other three sub plots with appropriate color maps (look up the cmap argument for `pyplot.imshow`).

- Remove the pixel coordinates from the axes as they are not relevant.

- Add the plot to your report.

15.2.3   Creating a gray scale image

You will perform that separation and, later, create a gray scale image using the following equation[2]

$$Y = 0.299R + 0.587G + 0.114B,$$

where R is the red channel value, G the green channel value, and B the blue channel value.

- Create a function `grayscale` that takes one argument (the filename of the image to make a grayscale version of).

- Load the image into an array.

- Gray scale images have the same value for each of the red, green and blue channels. You can create a single array for a gray scale image and use an appropriate color map to show it as a gray scale image. The other way is to create identical red, green, and blue channels and use `pyplot.imshow` to show that three channel image.

- Add the grayscale image to your report.

15.2.4   Detecting edges using a Sobel filter

In this section you will create a Sobel filter to detect edges in the image. Note you are not allowed to use `scipy.ndimage.filters.sobel` for this part of the assignment. Read the Wikipedia page on Sobel filters[3]. *This part the exercise is a bit harder than the earlier parts.*

- Create a new function `sobel_filter` that take a single argument (the filename of the file to operate on).

- Create a two-by-two grid of subplots (using `pyplot.sublots`).

- You can use array manipulations combined with the `scipy.signal.convolve` to implement the Sobel filter.

- Show the original image at the top left. The x-gradient on the top right, the y-gradient on the bottom left and the result of the complete Sobel filter at the bottom right.

- Make sure to properly label these plots.

- Add the resulting image to your report.

---

[2]With Y the luminance see `https://en.wikipedia.org/wiki/Grayscale`
[3]See `https://en.wikipedia.org/wiki/Sobel_operator`.

15.3   Hints

- If you are having problems with plots getting drawn on top of each other use the `pyplot.clf` function to clear the older plots.

Chapter 16

Kepler exoplanets (Group A, Week 6)

For this assignment you will use data from the Kepler telescope and the NASA exoplanet archive[1]. With the exoplanet archive data you can show that Kepler's third law holds.
Goals for this week:

- demonstrate the ability to design and implement simple programs,

- demonstrate the ability to collect and access tabular data (in text formats),

- demonstrate the ability to apply functions from Scipy to perform some simple analysis (fitting in this case).

Note: properly label any plots you make for this week.

16.1   Kepler's third law

We will start by showing that Kepler's third law holds using data from the NASA exoplanet archive.

- Go to the web pages of the NASA exoplanet archive and see what data are available. These webpages allow you to download data in simple formats like CSV or tab separated text files.

- Download the known exoplanet data in a format that is convenient.

- For the last few exercises you wrote your own file parsing code, you are now allowed to library functionality to make this easier. You could use the built-in `csv` module, you could use Numpy's `numpy.loadtxt` or `numpy.genfromtext`, or write your own parser for your data file. *Be sure to include all data files you use in your submission.*

- Create a scatter plot of the semi-major axis vs. orbital period. Add to this plot a line with the expected relation. Use logarithmic scaling of the axes as appropriate.

- Scipy offers functionality to perform fits, the functions that allow you to perform fits are available in the `scipy.optimize`. Read through that module's documentation (available from the Scipy website).

- Write a function perform a fit to the semi-major axis vs. orbital period data. Confirm that the relation as given by Kepler's third law holds. Create a new semi-major axis vs. orbital period scatter plot add to this figure a line plot based on the results of your fit. Color code the data points in your plot by the discovery method, make sure to add a legend explaining the colors.

- Based on the planetary orbits in the solar system confirm Keplers third law. Create a new scatter plot + fit to show the data and your fit. (Create a data file, do not hard code the planets and their parameters in your program.)

---

[1]See `http://exoplanetarchive.ipac.caltech.edu/`

- Find data about asteroid orbits, use that data to confirm Kepler's third law. Use at least 20 asteroid obits (that are not all part of one of the same asteroid belt). Add the names of all the ateroids you used to your report. Create a new scatter plot + fit to show the data and your fit.

- Write a small report explaining your methods and data sources.

- Create a plot of all exoplanets on the sky using the Kapteijn library[2] for Python, use the Hammer-Aitoff projection. If you cannot figure out how to use that Library, you are allowed to make a scatter plot of these sky positions using just Matplotlib (make sure to properly label your axes).

16.2   Handing in

Check that the following things are included in your submission:

- a report describing your data sources and methods,

- all code you wrote to perform your data analysis,

- the plots you were asked to create,

- any ancillary files needed to run your program.

Furthermore, check that your code conforms to proper style and that it has a clear structure. Check that your code does not depend on hardcoded paths, as those will not be present on the computer used for grading — make everything work with relative paths (and structure the tarball that you hand in such that the program will run without having to move files around).

---

[2]See `http://www.astro.rug.nl/software/kapteyn/allsky.html`

Chapter 17

Kepler exoplanets (Group B, Week 6)

For this assignment you will use data from the Kepler telescope and the NASA exoplanet archive[1]. With the Kepler data you can redetect previously discovered planets.
Goals for this week:

- demonstrate an ability to design and implement simple programs,

- demonstrate an ability to collect and access tabular data (in text formats),

- access data in FITS files, and run some data analysis on them (fitting and creating periodograms in this).

Note: properly label any plots you make for this week.

17.1   Redecting planets in Kepler data

For this assignment you will redetect previously known exoplanets in the Kepler data.

- Go to the web pages of the NASA exoplanet archive and see what data are available. These webpages allow you to download data in simple formats like CSV or tab separated text files. Furthermore the exoplanet archive also contains light curves observed by the NASA Kepler mission - that latter data is what you will need for this assignment.

- The Kepler mission observed a specific path of sky (the "Kepler field") for several years, repeatedly measuring the magnitudes of the stars in that field. The time-series created by Kepler show exoplanet transits, and by tracking these over a long time the orbits of these exoplanets can be reconstructed. Each of the stars in the Kepler field has a Kepler ID (sometimes called KOI for Kepler Object of Interest) that can be used to retrieve the data files, these IDs should not be confused with the numbering of the Kepler discoveries (given to planets like Kepler 5b).

- Find the earliest Kepler exoplanet discoveries (early discoveries have easily detected signatures in the data). The data can be downloaded from the Mikulski Archive for Space Telescopes (MAST)[2]. Choose one of these early discoveries and download a lightcurve data file that has a baseline that does not jump around to much (for convenience), and that shows a number of planet transits (the dips). Note: the website can generate plots for these data files so that you can check them before downloading.

- Kepler data files are in the FITS (Flexible Image Transport System) format that is used often in Astronomy. The Astropy library supports opening these files through PyFITS. If you have a recent version of PyFITS it will be part of the Astropy library otherwise it may be a standalone library. In the former case you need the `astropy.io.fits` module in the latter you should import `pyfits`. FITS files consists of a header(s) and data.

---

[1] See `http://exoplanetarchive.ipac.caltech.edu/`
[2] See `https://archive.stsci.edu/kepler/`.

- Write a function to open the FITS file (through PyFITS) and extract the raw time-series data and not the baseline corrected one — because part of the assignment is correct the baseline of the data. Create a plot of the photometry data, make sure your time axis is correct.

- Inspect your time-series and choose a range of data that has no breaks, and filter out any not-a-number values as they will cause problems in the following analysis.

- Scipy provides a way of creating Lomb Scargle periodograms, we will use that functionality (do not use the Lomb Scargle functions from other libraries like AstroML).

- Scipy's implementation of these periodograms expects data that is centered around zero. The Kepler data is not like that. Fit a low order polynomial function to the data (so that the planet transits are not fit) and subtract this from the data. You will now have data that is close to centered around zero. Furthermore subtract the median of the data to get closer to zero (while ignoring the planet transits).

- Use Scipy's periodogram functionality to find the period of the planet's orbit. Create a plot of the periodogram and indicate where the planets signal can be found.

- Extend your program such that it can process several lightcurve data files at once. Show that it works for several planets and several lightcurves per planet.

- Write a small report explaining your methods and data sources.

17.2   An all sky plot

In this part of the assignment you will create a plot of the positions on the sky of the known exoplanets.

- The NASA exoplanet archive also provides data on all known exoplanets. This data is available as simple text files like the ones you processed before.

- Download the sky positions of the known exoplanets, and write a parser for these data. You are allowed to use a library function to do the parsing (either from csv, Numpy or Scipy). *Be sure to include all data files you use in your submission.*

- Create a plot of all exoplanets on the sky using the Kapteijn library[3] for Python, use the Hammer-Aitoff projection. If you cannot figure out how to use that Library, you are allowed to make a scatter plot of these sky positions using just Matplotlib (make sure to properly label your axes).

17.3   Handing in

Check that the following things are included in your submission:

- a report describing your data sources and methods,

- all code you wrote to perform your data analysis,

- the four plots you were asked to create,

- any ancillary files needed to run your program.

Furthermore, check that your code conforms to proper style and that it has a clear structure. Check that your code does not depend on hardcoded paths, as those will not be present on the computer used for grading — make everything work with relative paths (and structure the tarball that you hand in such that the program will run without having to move files around).

---

[3]See http://www.astro.rug.nl/software/kapteyn/allsky.html

Chapter 18

Some analysis on a globular cluster (Week 7 groups A and B)

This week's assignment contains some harder parts, so the grading will be as follows (you can score a total of 9 points):

- 3 points for style and program structure — graded strictly

- 4 points for the plotting of FITS files

- 1 point for the detecting stars part of the exercise

- 1 point for the Globular cluster radial stellar density profile

As always, take care to properly label your plots. Also, create a small report with these plots and an explanation of your program.

18.0.1 Plotting FITS files

You will furthermore need Astropy to access FITS data files.

- Download the NGC884 images from the `astro.mprog.nl` website, make sure you get all of them.

- Use PyFits, which nowadays is part of Astropy, to open the FITS files. You will need the `astropy.io.fits` module and its `open` function that can load FITS files.

- FITS files consist of one or more tables (Header Data Units) that each have metadata in a header and the data itself. Access the header and answer the following questions in your report: what is the size of the images (in pixels), what is the focal length of the telescope and what is the size of its mirror.

- Use `pyplot.imshow` to show the images, you may want to select only part of the image (it is quite large, meaning that the stars in it may be hardly noticeable when you show the full image). You should choose an appropriate colormap, show a color bar alongside the image and change the colormap so that it uses a logarithmic scaling. That logarithmic scaling will show more stars.

- Create an image showing 4 subplots, one for the green filter, one for the red filter, one for the blue filter and one showing them as one image in full color. The scalings should be logarithmic but you are free to adjust the ranges you show to taste. If the images are misaligned slightly do not worry about that.

18.0.2 Detecting stars

For this part of the exercise you will need the Photutils library[1] which provides functionality for photometry.

- Check that you have photutils installed, if not please install it.

- In the documentation for `photuils` read the page about source detection. Choose one of the techniques described on that page and run it on one of the NGC884 images. Create an image combining the FITS file image with a scatter plot that shows the detected stellar positions. Do not use the `astropy.visualization` module - you should be able to create this plot using Matplotlib directly.

- Create two one dimensional histograms, one showing the distribution of stars along the horizontal axis, and one along the vertical axis.

18.0.3 Globular cluster M5 in SDSS

The Sloan Digital Sky Survey (SDSS) observed a large part of the sky and all its data is available online. We will use the M5 globular cluster for this assignment as it is available in the SDSS data. The goal of this part of the assignment is to create a histogram showing the number of stars as a function of distance from the cluster center.

- Go to the bulk download page[2] of SDSS and search for the M5 data and download at the data in the different filters. Use the U-band image that is available.

- Detect the stars as before.

- Create an image showing the cluster image in the background and overplot a scatter plot of the star positions (assume all detected sources are in fact stars).

- Find the cluster center's x and y pixel coordinates in the image. You can do this for example by creating a histogram of the source X positions and fitting a gaussian to that histogram. The `pyplot.hist` function returns the histogram values, the edges of the histogram bins and somethings called patches (used by `pyplot` to do the actual drawing — of no interest for now). You need the bin centers (not the edges) for this step. The gaussian fit gives you the X coordinate of the cluster center. Repeat these steps for the Y cluster center in pixel coordinates.

- Go through the detected sources, calculate the distance to the cluster center and then create a histogram of these distances. You can turn the pixel coordinates into physical units (angle on the sky) with information from the FITS file header — but for the purposes of this part of the exercise pixel distances are enough.

---

[1]See https://photutils.readthedocs.org/en/latest/.
[2]See http://data.sdss3.org/bulkFields/names.

Chapter 19

Software licensing

While software licensing is not a technical subject, it is my opinion that every programmer should have a basic understanding of licensing, so that he or she can make informed decisions in these matters. Much like in research ideas of others should be credited properly, you should respect the licenses that software comes under. In the context of software a license describes what that software can be used for and under what circumstances copies of the software can be made legally. Every piece of software that is created is automatically protected by copyright, that is the right to make copies of it lies with its author (and *only* with its author).

One of the reasons why Python and its "eco system" is so popular is because it can freely be copied, shared and extended by anyone. This is possible because Python and most libraries written for Python are released under licenses that allow use for any purpose, copying and extension. The license used by the Python progject, the so-called Python Software License, is an example of a larger group of Free Software or Open Source licenses. In this chapter I will explain some of the basic ideas behind these licenses and the implications of a particular choice of license.

19.1   The origin of Free and Open Source Software

Early computer hardware usually came with software that was freely shared. In the 1970s increasingly software itself was sold commercially. A culture of sharing and adapting software to the users' needs had existed before. In the 1980s Richard Stallman started work on reimplementation of UNIX whose source would be freely available, he also started the Free Software Foundation to support this effort. Stallman insisted that people have rights when it comes to software, the modern definition of software freedoms enumerates 4 rights:

- The freedom to run the program, for any purpose (freedom 0).

- The freedom to study how the program works, and change it so it does your computing as you wish (freedom 1). Access to the source code is a precondition for this (freedom 1).

- The freedom to redistribute copies so you can help your neighbor (freedom 2).

- The freedom to distribute copies of your modified versions to others (freedom 3). By doing this you can give the whole community a chance to benefit from your changes. Access to the source code is a precondition for this.

To *redistribe* in this context means to give or sell someone else a copy of the software. To make sure that their software would stay free the FSF created a special license called the General Public License (GPL). It governs the act of (re-)distributing the software, a user who receives software under GPL can use it for any purpose, can request the source, is allowed to change that source and provide others with copies. What the GPL does not allow is to put further restrictions on the use of the piece of software. One particularly clever aspect of the GPL is that it uses the rules of copyright to guarantee extra rights for its users; when you create copies of software under GPL you either abide by the rules of GPL or you have no right whatsoever to make copies as per the regular rules of copyright.

This aspect is said to make the GPL a *copyleft* license. The GPL has stood the test of time and is a quite popular license nowadays, e.g. Linux uses GPL version 2.

Open source was created by some who thought Free Software was too politcal and wanted a more pragmatic definition. The source to open source software is available, but more restrictions may apply when compared to the Free Software definition. Free software licenses are a subset of Open source licenses. Which licenses can be called open source is governed by the Open Source Initiative (OSI). Toghether Free Software and Open Source software are also referred to as Free / Open Source Software (FOSS).

19.2 Popular licenses

In this section a few popular FOSS licenses are described, necessarily in a condensed fashion. More complete information about FOSS licenses in general can be found on the website of the FSF[1].

19.2.1 GPL, LGPL and AGPL

Software licensed under GPL can be used for any purpose, its source must be available and any modifications to the source of the software must be shared upon redistribution. This implies that when you build software using GPL source, your modifications or additions are also bound by the GPL. This last property may not always be desireable when you are for instance building a library that is to be used generally. The LGPL (*L* for *lesser* or *library*) allows you to link your software to a LGPL library without having to share the source the source to your software. It however does require you, upon redistribution, to share changes to the LGPL library itself.

An implication of building software using GPL sources is that you cannot restrict what others that receive the software are allowed to do with it. As an example if you are working on some GPL software that you share with a fellow researcher, but don't want to share with the whole world yet, you cannot require the recipient to keep the software private.

The GPL has a loophole when it comes to online services. Since GPL covers the distribution of software it does not extend any rights to users of that service. So an online service may be built on modified GPL code without sharing it, in effect complying with the letter of the license while breaking it in spirit. The AGPL (*A* for Affero) was created to close this loophole, services built using AGPL code must offer users of that service the possibility to download the underlying source code.

19.2.2 The BSD, MIT and X licenses

The Berkely Software Distribution (BSD) licenses, the Massachussets Institute of Technology (MIT) and X licenses are more permissive than the *GPL licenses. They are not copleft, i.e. they do not require recipients of the software and source to share their modifications and allow further restrictions to be placed on the software. There are two versions of the BSD license you may come across the old, two clause, and the new, three clause, one. The former cannot be combined with GPL source because it requires the software (and not just the source) to display the original authors somewhere, which is an extra restriction on top of what the GPL would require. You should use the newer license. If you are writing a new piece of software and you want recipients of that software to share their modifications to it the BSD and MIT licenses may not be a good choice.

---

[1]The FSF website can be found at `http://www.fsf.org`.

Chapter 20

Text editors

- Kate and Gedit are the editors included with, respectively, the KDE and Gnome desktop environments for Linux. Both are graphical editors that support syntax highlighting in many programming languages. Kate, like all things KDE, is very configurable. (Recommended.)

- Xcode is Apple's integrated development environment for Mac OS X. While it supports Python, it includes many features that are not needed for Python. (Avoid, unless you use it for other languages already.)

- Textmate is a capable text editor for Mac OS X, previously one of the favorites for that system. Version 1.x is available commercially while version 2 is available as open source software. (Recommended.)

- Sublime is a modern grahical text editor for Mac OS X that supports many programming languages. (recommended).

- Notepad++ is a nice programming text editor for Windows that is completely free to use. (Recommended)

- IDLE is the editor and interactive Python shell combination included with Python by default. IDLE is cross platform, not integrating well with some of its hosts. The interactive shell can be very slow when your program prints a large amount of text. IDLE may also complicate debugging because at times you will come across tracebacks that point to IDLE's code. Since IDLE only supports Python syntax highlighting you will end up learning another editor in the long run. (Avoid, unless required by some course.)

- Emacs is a very capable and customizable editor. It was originally a text based editor although there are also graphical versions like XEmacs. The learning curve for Emacs is steeper than graphical editors like Kate or Gedit. Emacs will also work in a terminal connected to a remote computer system, this allows you to do programming work on a remote system that does not allow graphical connections. (Recommended.)

- vi or vim are very capable and customizable editors that have a *very steep learning curve*. Vi(m) uses very concise commands that can be combined to perform complex tasks. The vim, "vi improved", editor has more features than vi. Like Emacs, vi(m) will work in a terminal and allow you to do programming on systems without graphical displays. (Recommended, if you have the time to learn it.)

- pico and nano are simple text based editors one of which is usually available on a remote Linux systems. Use these editors only for quick small edits on remote systems, and only if you cannot use a more powerful editor like emacs and vi(m). (Avoid for anything other than emergencies.)